

# Out-of-Order Execution & Dynamic Scheduling

---

15-740 FALL'18

NATHAN BECKMANN

# So Far ...

---

$$\begin{aligned}\text{CPU Time} &= \text{CPU clock cycles} \times \text{clock cycle time} \\ &= \frac{\text{instructions}}{\text{program}} \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}}\end{aligned}$$

Increase clock frequency

Decrease CPI

- Pipeline:
  - divide CPI by # of stages + bubbles
  - add bypass
- Superscalar
  - divide CPI by issue width + bubbles from
    - data dependencies
    - control dependencies

# What next?

---

What slows us down here?

```
// for (i=0; i<N; i++) c[i] = a[i]*b[i]+i*2;
```

```
loop: ld      r1, (r2+r9*8)
      ld      r3, (r4+r9*8)
      mult    r5, r3, r1
      mult    r3, r9, 2
      add     r5, r5, r3
      st      r5, (r6+r9*8)
      add     r9, r9, 1
      cmp     r9, r10
      bnz     loop
```

# What next?

---

What slows us down here?

```
// for (i=0; i<N; i++) c[i] = a[i]*b[i]+i*2;
```

```
loop: ld      r1, (r2+r9*8)
      ld      r3, (r4+r9*8)
      mult    r5, r3, r1
      mult    r3, r9, 2
      add     r5, r5, r3
      st      r5, (r6+r9*8)
      add     r9, r9, 1
      cmp     r9, r10
      bnz     loop
```

# How can we improve this?

---

Assume:

- 2-issue superscalar
- 1 LD/ST unit (2 cycles),
- 1 mult (2 cycles)
- 1 add (1 cycle)

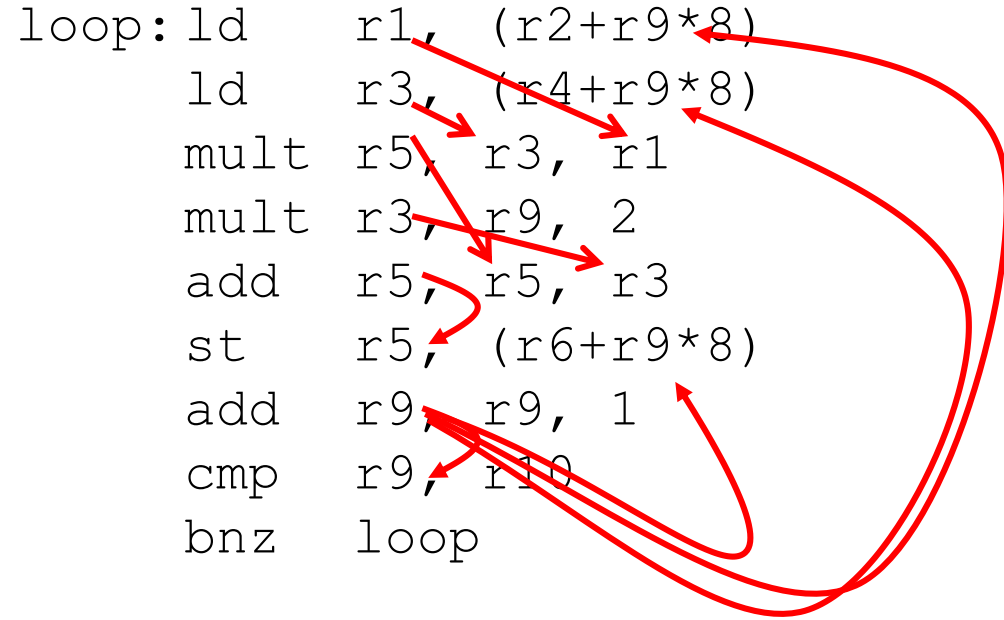
```
ld      r1, (r2+r9*8)
..ld     r3, (r4+r9*8)
..mult   r5, r3, r1
..mult   r3, r9, 2
..add    r5, r5, r3
.st      r5, (r6+r9*8)
.add     r9, r9, 1
.cmp     r9, r10
.bnz     loop
```

12 cycles

# What next?

What slows us down here?

```
loop: ld    r1, (r2+r9*8)
      ld    r3, (r4+r9*8)
      mult  r5, r3, r1
      mult  r3, r9, 2
      add   r5, r5, r3
      st    r5, (r6+r9*8)
      add   r9, r9, 1
      cmp   r9, r10
      bnz   loop
```

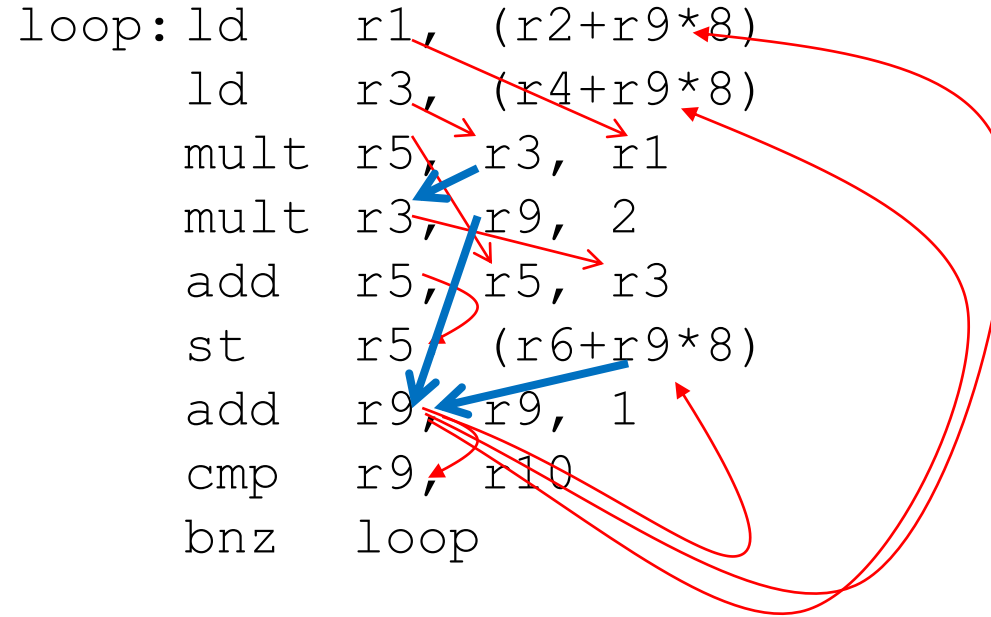


## Data Dependencies

- RAW: Read after write. (true dependence)

# What next?

What slows us down here?

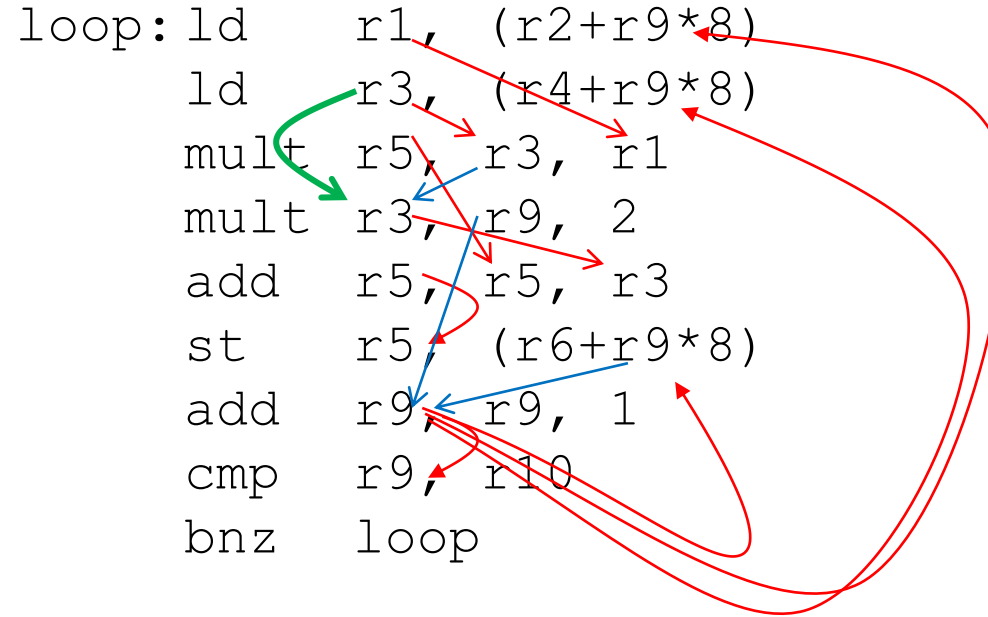


## Data Dependencies

- RAW: Read after write. (true dependence)
- WAR: write after read (false dependence)

# What next?

What slows us down here?



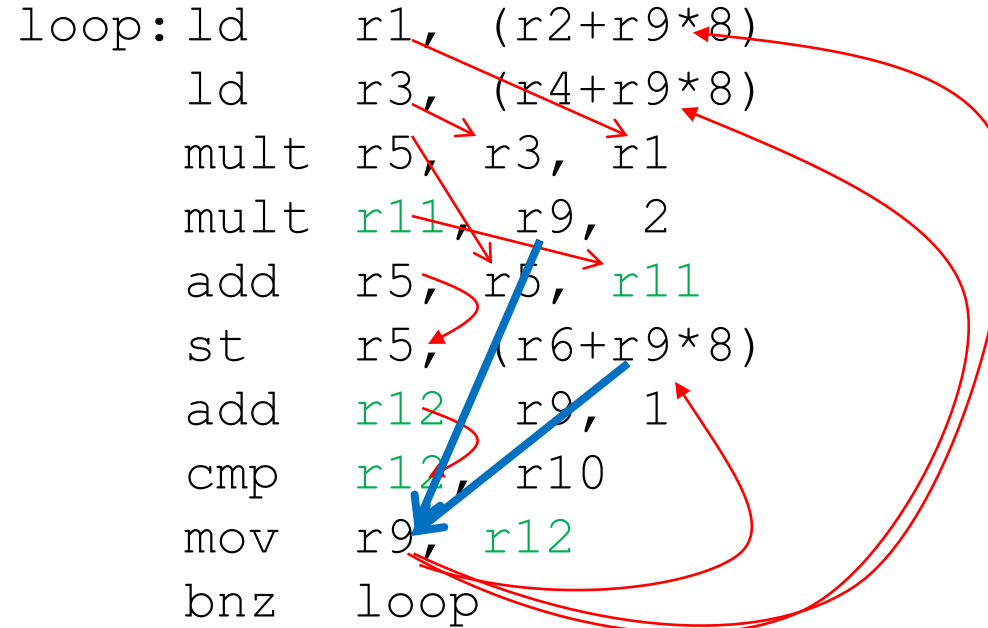
## Data Dependencies

- RAW: Read after write. (true dependence)
- WAR: write after read (false dependence)
- WAW: write after write (false dependence)



# Eliminating WAR (and WAW)

What slows us down here?



## Data Dependencies

- RAW: Read after write. (true dependence)
- WAR: write after read (false dependence)
- WAW: write after write (false dependence)

# Reschedule for better ILP

---

```
loop: ld    r1, (r2+r9*8)
      ld    r3, (r4+r9*8)
      mult  r5, r3, r1
      mult  r11, r9, 2
      add   r5, r5, r11
      st    r5, (r6+r9*8)
      add   r12, r9, 1
      cmp   r12, r10
      mov   r9, r12
      bnz   loop
```

# Can we do better?

---

```
loop: ld      r1, (r2+r9*8)
      mult    r11, r9, 2
      ..ld    r3, (r4+r9*8)
      ..mult   r5, r3, r1
      ..add    r5, r5, r11
      .st      r5, (r6+r9*8)
      add     r12, r9, 1
      .cmp    r12, r10
      mov     r9, r12
      bnz     loop
```


10 cycles

# More Rescheduling

---

```
loop: ld    r1, (r2+r9*8)
      mult  r11, r9, 2
      ld    r3, (r4+r9*8)
      mult  r5, r3, r1
      add   r5, r5, r11
      st    r5, (r6+r9*8)
      add   r12, r9, 1
      cmp   r12, r10
      mov   r9, r12
      bnz   loop
```

```
loop: ld    r1, (r2+r9*8)
      mult  r11, r9, 2
      ld    r3, (r4+r9*8)
      add   r12, r9, 1
      cmp   r12, r10
      mult  r5, r3, r1
      add   r5, r5, r11
      st    r5, (r6+r9*8)
      mov   r9, r12
      bnz   loop
```



# More Rescheduling

---

```
ld      r1, (r2+r9*8)
mult    r11, r9, 2
..ld     r3, (r4+r9*8)
add     r12, r9, 1
.cmp    r12, r10
mult    r5, r3, r1
..add   r5, r5, r11
.st     r5, (r6+r9*8)
mov     r9, r12
.bnz    loop
```

7 cycles

# What Is Holding Us Back?

---

```
ld      r1, (r2+r9*8)
mult    r11, r9, 2
..ld     r3, (r4+r9*8)
add     r12, r9, 1
.cmp     r12, r10
mult    r5, r3, r1
..add    r5, r5, r11
.st      r5, (r6+r9*8)
mov     r9, r12
.bnz    loop
```

- In order issue
- Not enough Function units
- Function units too slow
- Not enough register names
- Control dependence
- Static scheduling

# What Is Holding Us Back?

---

```
ld      r1, (r2+r9*8)
mult    r11, r9, 2
..ld     r3, (r4+r9*8)
add     r12, r9, 1
.cmp     r12, r10
mult    r5, r3, r1
..add    r5, r5, r11
.st      r5, (r6+r9*8)
bnz     loop
ld      r1, (r2+r12*8)
```

- In order issue
- Not enough Function units
- Function units too slow
- **Not enough register names**
- Control dependence
- Static scheduling

# What Is Holding Us Back?

---

- In order issue
- Not enough Function units
- Function units too slow
- Not enough register names
- Control dependence
- Static scheduling

```
ld      r1, (r2+r9*8)
..ld     r3, (r4+r9*8)
..mult   r5, r3, r1
..mult   r3, r9, 2
..add    r5, r5, r3
.st      r5, (r6+r9*8)
.add     r9, r9, 1
.cmp     r9, r10
.bnz     loop
```

Hardware

```
ld      r1, (r2+r9*8)
mult    r11, r9, 2
..ld     r3, (r4+r9*8)
add     r12, r9, 1
.cmp     r12, r10
mult    r5, r3, r1
..add    r5, r5, r11
.st      r5, (r6+r9*8)
bnz     loop
```



# What Is Holding Us Back?

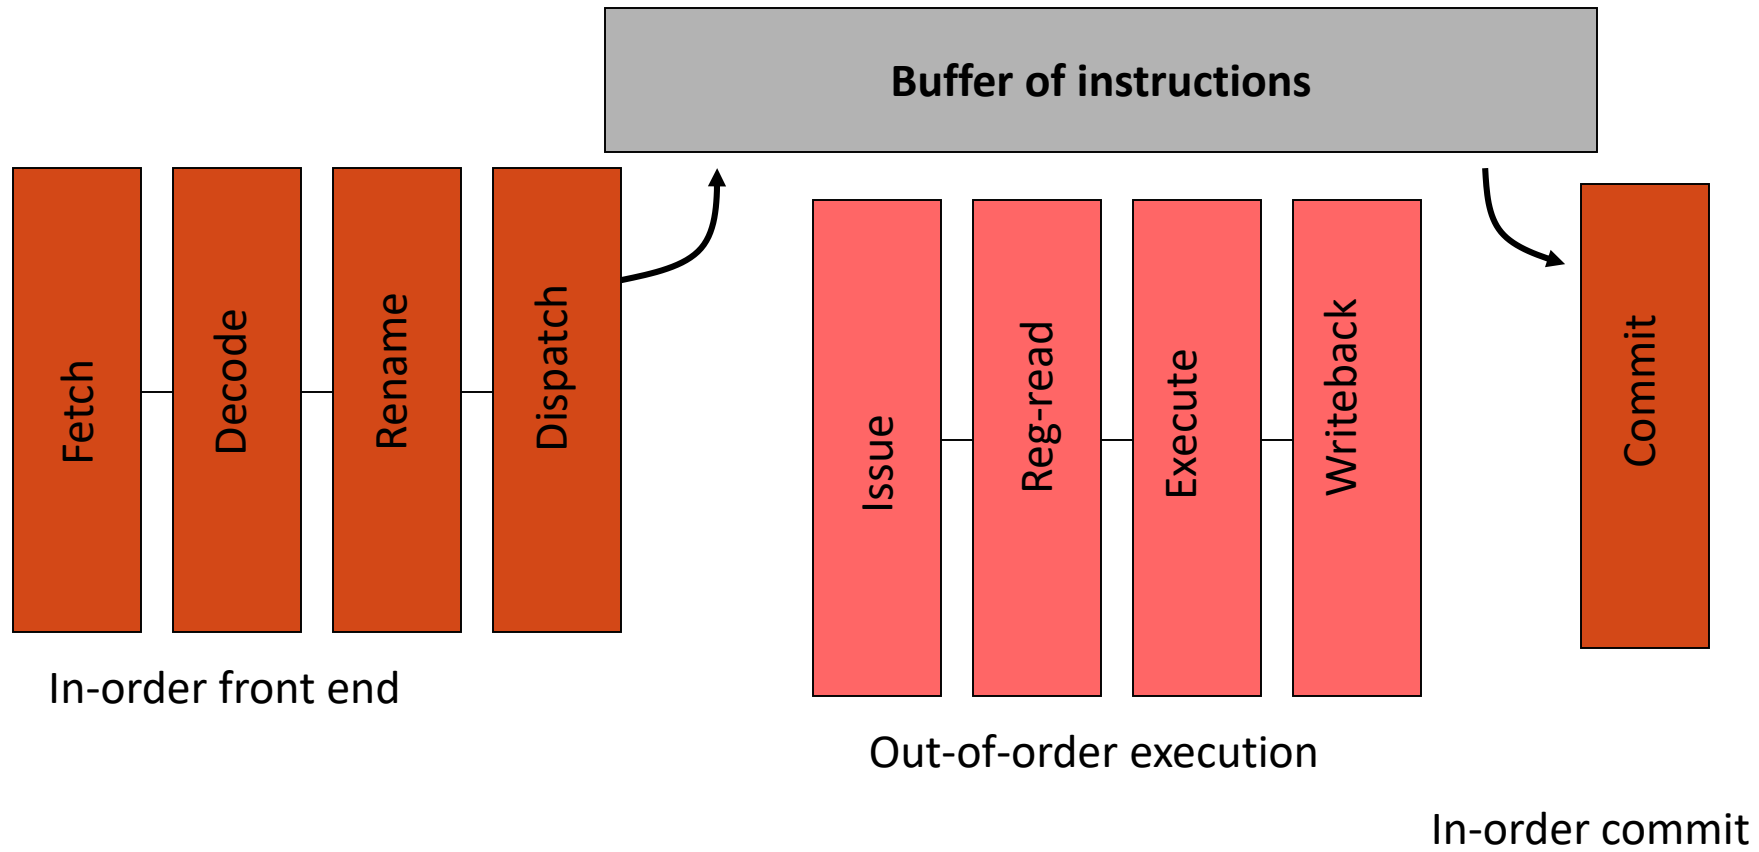
---

```
ld      r1, (r2+r9*8)
mult    r11, r9, 2
..ld     r3, (r4+r9*8)
add     r12, r9, 1
.cmp     r12, r10
mult    r5, r3, r1
..add    r5, r5, r11
ld      r1, (r2+r12*8)
.st     r5, (r6+r9*8)
bnz     loop
```

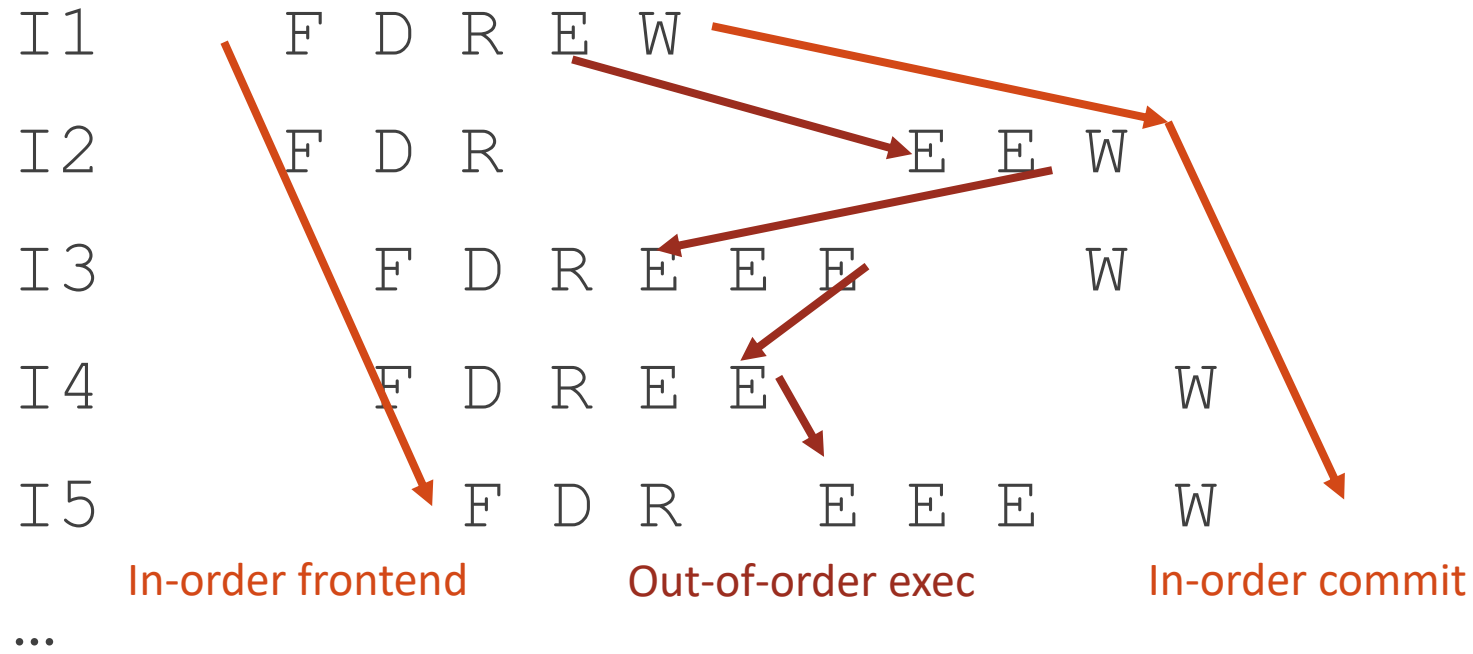
- In order issue
- Not enough Function units
- Function units too slow
- Not enough register names
- **Control dependence**
- Static scheduling

# Out-of-Order Pipeline

---



# Out-of-Order Execution



# Out-of-Order Execution

---

Also called “Dynamic scheduling”

- Done by the hardware on-the-fly during execution

Looks at a “window” of instructions waiting to execute

- Each cycle, picks the next ready instruction(s)

Two steps to enable out-of-order execution:

Step #1: Register renaming – to avoid “false” dependencies

Step #2: Dynamically schedule – to enforce “true” dependencies

Key to understanding out-of-order execution:

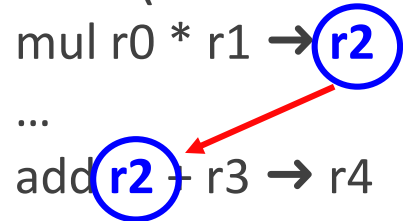
- **Data dependencies**

# Dependence types

---

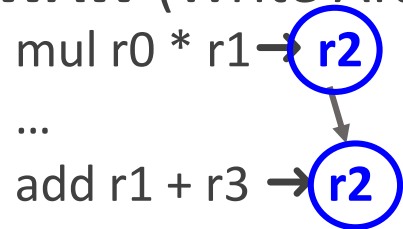
**RAW** (Read After Write) = “true dependence” (true)

mul r0 \* r1 → **r2**  
...  
add **r2** + r3 → r4



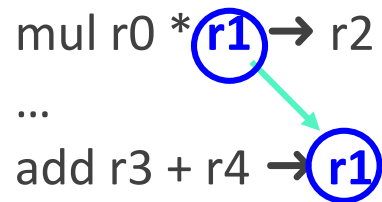
**WAW** (Write After Write) = “output dependence” (false)

mul r0 \* r1 → **r2**  
...  
add r1 + r3 → **r2**



**WAR** (Write After Read) = “anti-dependence” (false)

mul r0 \* **r1** → r2  
...  
add r3 + r4 → **r1**



*WAW & WAR are “false”, Can be **totally eliminated** by “renaming”*

# Register Renaming

---

# Step #1: Register Renaming

To eliminate register conflicts/hazards

“Architected” vs “Physical” registers – level of indirection

- Names: **r1**, **r2**, **r3**
- Locations: **p1**, **p2**, **p3**, **p4**, **p5**, **p6**, **p7**
- Original mapping: **r1**→**p1**, **r2**→**p2**, **r3**→**p3**, **p4**–**p7** are “available”

MapTable

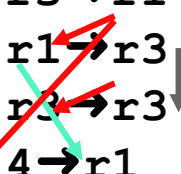
r1	r2	r3
p1	p2	p3
p4	p2	p3
p4	p2	p5
p4	p2	p6

FreeList

p4, p5, p6, p7
p5, p6, p7
p6, p7
p7

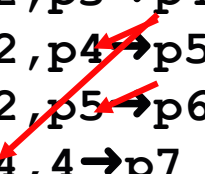
Original insns

```
add r2, r3 → r1
sub r2, r1 → r3
mul r2, r3 → r3
div r1, 4 → r1
```



Renamed insns

```
add p2, p3 → p4
sub p2, p4 → p5
mul p2, p5 → p6
div p4, 4 → p7
```



- Renaming – conceptually write each register once
  - + Removes **false** dependences
  - + Leaves **true** dependences intact!
- When to reuse a physical register? **After overwriting insn done**

# Register Renaming Algorithm (Simplified)

---

Two key data structures:

- `maptable[architectural_reg] ➔ physical_reg`
- Free list: allocate (new) & free registers (implemented as a queue)

Algorithm: at “decode” stage for each instruction:

```
insn.phys_input1 = maptable[insn.arch_input1]
```

```
insn.phys_input2 = maptable[insn.arch_input2]
```

```
new_reg = new_phys_reg()
```

```
maptable[insn.arch_output] = new_reg
```

```
insn.phys_output = new_reg
```



# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3      →      xor p1 ^ p2 →  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3      →      xor p1 ^ p2 → p6  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3      →      xor p1 ^ p2 → p6  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table


p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 →

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table


p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table


p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table


p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 →

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10


Free-list



# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table


p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 →

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

# Renaming example


---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

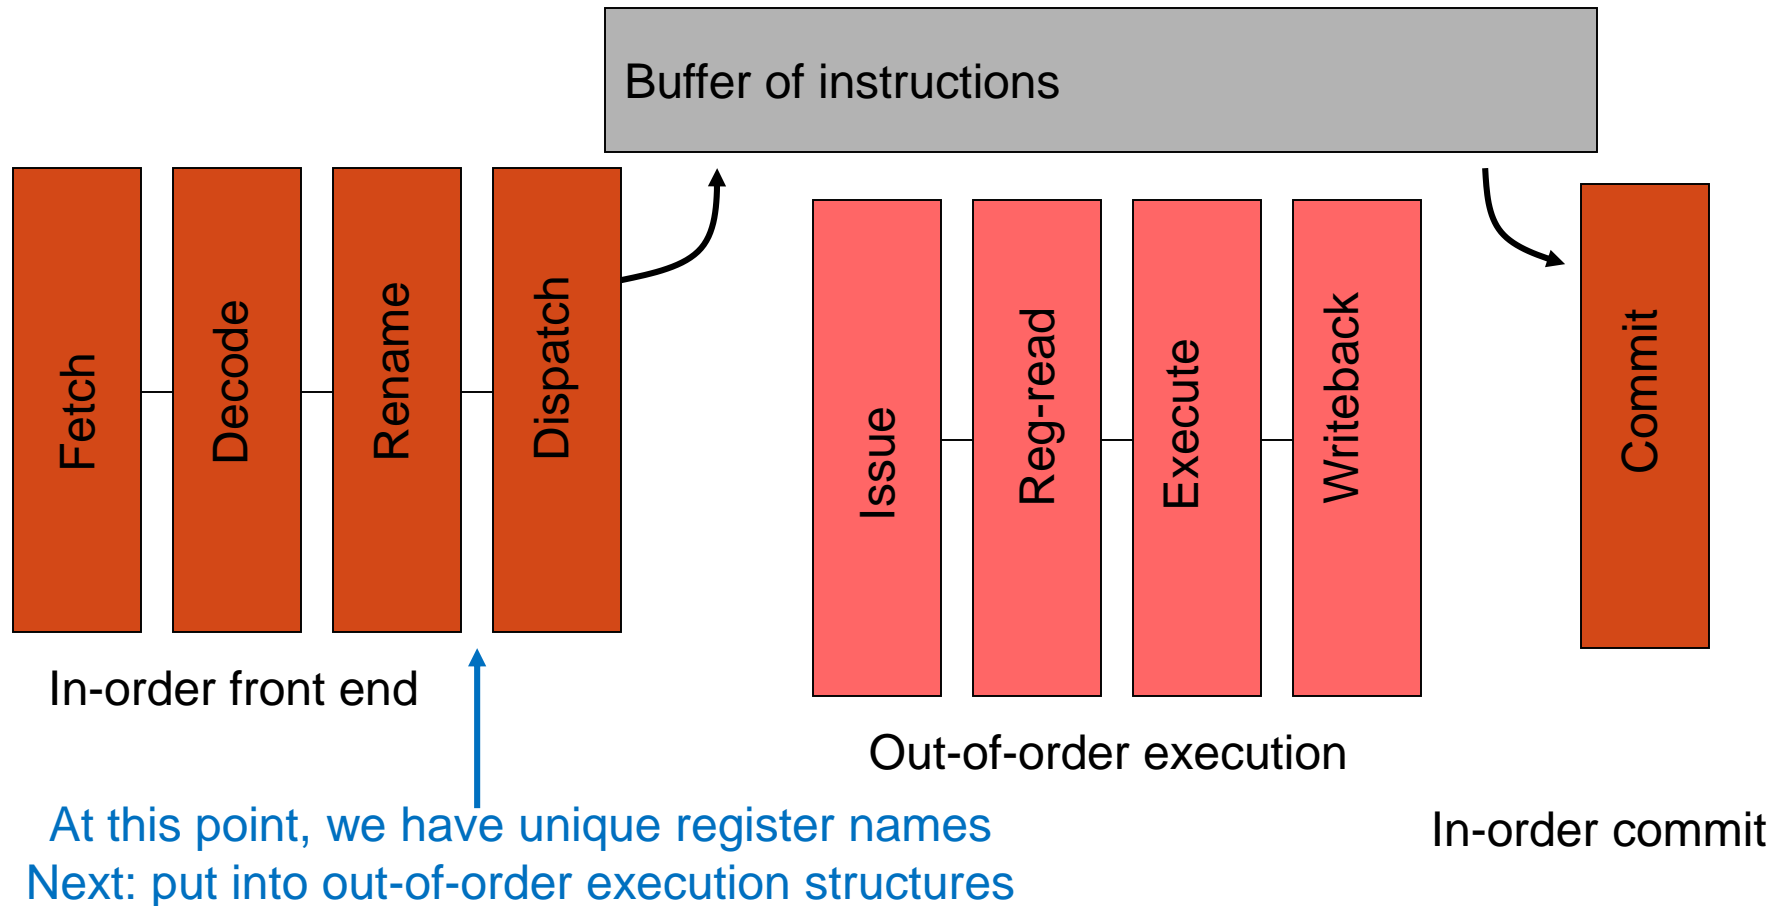


Free-list


# Dynamic Scheduling Mechanisms

---

# Out-of-order Pipeline



# Step #2: Dynamic Scheduling



Ready Table

	P2	P3	P4	P5	P6	P7
Time	Yes	Yes				
	Yes	Yes	Yes			
	Yes	Yes	Yes	Yes		Yes
	Yes	Yes	Yes	Yes	Yes	Yes

add p2, p3 → p4  
sub p2, p4 → p5 and div p4, 4 → p7  
mul p2, p5 → p6

Instructions fetch/decoded/renamed into *re-order buffer (ROB)*

Check which instructions are ready and execute earliest ready instruction



# Out-of-order pipeline


---

Execution (out-of-order) stages

**Select** ready instructions

- Send for execution

**Wakeup** dependents



# Dispatch

---


Put renamed instructions into out-of-order structures

## Re-order buffer (ROB)

- Holds instructions until commit
- Preserves in-order commit

## Issue Queue

- Central piece of scheduling logic
- Holds un-executed instructions
- Tracks ready inputs
  - Physical register names + ready bit
  - “AND” the bits to tell if ready



# Dispatch Steps

---

Allocate Issue Queue (IQ) slot

- Full? **Stall!**

Read **ready bits** of inputs

- 1-bit per physical reg

Clear **ready bit** of output in table

- Instruction has not produced value yet

Write data into Issue Queue (IQ) slot

# Dispatch Example

---

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	y
p8	y
p9	y

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	y
p8	y
p9	y

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	y
p9	y

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	y

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

Issue Queue

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	n	---	y	p9	3

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	n



# Dynamic Scheduling/Issue Algorithm

---

Data structures:

- Ready table[phys\_reg] → yes/no (part of “issue queue”)

Algorithm at “issue” stage (prior to read registers):

```
foreach instruction:
```

```
    if table[insn.phys_input1] == ready &&  
        table[insn.phys_input2] == ready then  
        insn is “ready”
```

```
select the oldest “ready” instruction
```

```
    table[insn.phys_output] = ready
```

Multiple-cycle instructions?

- For an insn with latency of N, set “ready” bit N-1 cycles in future (like scoreboarding)

# Issue = Select + Wakeup

---

## Select oldest of “ready” instructions

- Single-issue: “xor” is the oldest ready instruction below
- Dual-issue: “xor” and “sub” are the two oldest ready instructions below
- Note: may have structural hazards (i.e., load/store/floating point)

Insn	Inp1	R	Inp2	R	Dst	Bday	
xor	p1	y	p2	y	p6	0	Ready!
add	p6	n	p4	y	p7	1	
sub	p5	y	p2	y	p8	2	Ready!
addi	p8	n	---	y	p9	3	

# Issue = Select + Wakeup

## Wakeup dependent instructions

- Search for destination (Dst) in inputs & set “ready” bit
  - Implemented with a special memory array circuit called a Content Addressable Memory (CAM)
- Also update ready-bit table for future instructions

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0
add	p6	y	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	y	---	y	p9	3

### Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	n
p8	y
p9	n

# Issue

---

**Select/Wakeup** one cycle


Dependent instructions execute on back-to-back cycles

- Next cycle: add/addi are ready:


Issued instructions are removed from issue queue

Insn	Inp1	R	Inp2	R	Dst	Bday
add	p6	y	p4	y	p7	1
addi	p8	y	---	y	p9	3


# OOO execution (2-wide)




# OOO execution (2-wide)



# OOO execution (2-wide)




# OOO execution (2-wide)






# OOO execution (2-wide)



# OOO execution (2-wide)

---



# When Does Register Read Occur?

---

Current approach: after select, right before execute

- **Not during in-order part of pipeline, in out-of-order part**
- Read **physical** register (renamed)
- Or get value via bypassing (based on physical register name)
- This is Pentium 4, MIPS R10k, Alpha 21264, IBM Power4, Intel's "Sandy Bridge" (2011)
- Physical register file may be large → Multi-cycle read

Older approach:

- Read as part of "issue" stage, keep values in Issue Queue / "Reservation Stations"
  - At commit, write them back to single "architectural register file"
- Pentium Pro, Core 2, Core i7
- Simpler, but may be less energy efficient (more data movement)

# Renaming Revisited

---

# Re-order Buffer (ROB)

---

ROB entry holds all info for recovery/commit

- **All instructions** & in order
- Architectural register names, physical register names, insn type
- Not removed until very last thing (“commit”)

Operation

- Dispatch: insert at tail (if full, stall)
- Commit: remove from head (if not yet done, stall)

Purpose: in-order commit

- Appearance of in-order execution
- Precise exceptions
- Enables speculation

# Renaming revisited

---

Track (or “log”) the “overwritten register” in ROB

- Free this register at commit
- Also used to restore the map table on “recovery”
  - Branch mis-prediction recovery

# Register Renaming Algorithm (Full)

---

Two key data structures:

- `maptable[architectural_reg] ➔ physical_reg`
- Free list: allocate (new) & free registers (implemented as a queue)

Algorithm: at “decode” stage for each instruction:

```
insn.phys_input1 = maptable[insn.arch_input1]
insn.phys_input2 = maptable[insn.arch_input2]
insn.old_phys_output = maptable[insn.arch_output]
new_reg = new_phys_reg()
maptable[insn.arch_output] = new_reg
insn.phys_output = new_reg
```

**At “commit”**

- **Once all older instructions have committed, free register**  
**`free_phys_reg(insn. old_phys_output)`**

# Recovery

---

## Completely remove wrong path instructions

- Flush from IQ
- Remove from ROB
- Restore map table to before misprediction
- Free destination registers

## How to restore map table?

- **Option #1:** log-based reverse renaming to recover each instruction
  - Tracks the old mapping to allow it to be reversed
  - Done sequentially for each instruction (slow)
- **Option #2:** checkpoint-based recovery
  - Checkpoint state of maptable and free list each cycle
  - Faster recovery, but requires more state
- **Option #3:** hybrid (checkpoint for branches, unwind for others)



# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

————→ xor p1 ^ p2 →

[ p3 ]

r1	p1
r2	p2
r3	<b>p3</b>
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

————→ xor p1 ^ p2 → p6

[ p3 ]

r1	p1
r2	p2
r3	<b>p6</b>
r4	p4
r5	p5

Map table


p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 →

[ p3 ]  
[ p4 ]

r1	p1
r2	p2
r3	p6
r4	<b>p4</b>
r5	p5

Map table


p7
p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7

[ p3 ]  
[ p4 ]

r1	p1
r2	p2
r3	p6
r4	<b>p7</b>
r5	p5

Map table


p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 →

[ p3 ]  
[ p4 ]  
[ p6 ]

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table


p8
p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8

[ p3 ]  
[ p4 ]  
[ p6 ]

r1	p1
r2	p2
r3	<b>p8</b>
r4	p7
r5	p5

Map table


p9
p10

Free-list

# Renaming example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 →

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list



# Renaming example

---


xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Recovery Example

Now, let's use this info. to recover from a branch misprediction

**bnz r1, loop**

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1


**bnz p1, loop**

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[   ]  
[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Recovery Example

---


bnz r1 loop  
xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

bnz p1, loop  
xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[   ]  
[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Recovery Example

---

bnz r1 loop  
xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3

bnz p1, loop  
xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → **p8**

[   ]  
[ p3 ]  
[ p4 ]  
**[ p6 ]**

r1	p1
r2	p2
r3	<b>p6</b>
r4	p7
r5	p5

Map table

<b>p8</b>
p9
p10

Free-list

# Recovery Example

---

bnz r1 loop  
xor r1 ^ r2 → r3  
add r3 + r4 → r4

bnz p1, loop  
xor p1 ^ p2 → p6  
add p6 + p4 → **p7**

[   ]  
[ p3 ]  
[ **p4** ]

r1	p1
r2	p2
r3	p6
r4	<b>p4</b>
r5	p5

Map table

<b>p7</b>
p8
p9
p10

Free-list

# Recovery Example

---

bnz r1 loop  
xor r1 ^ r2 → r3

bnz p1, loop  
xor p1 ^ p2 → p6

[   ]  
[ p3 ]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Recovery Example

---

bnz r1 loop

bnz p1, loop

[   ]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Commit

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

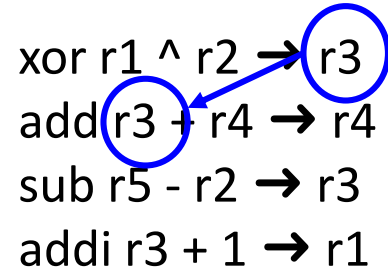
- Commit: instruction becomes **architected state**
  - In-order, only when instructions are finished
  - Free overwritten register (*why only at commit?*)



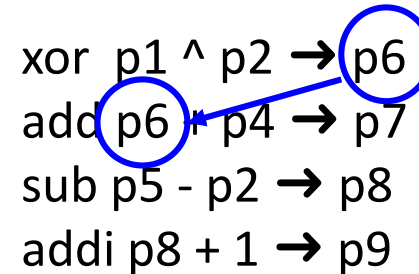
# Freeing over-written register

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9



[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

- P3 was r3 before xor
- P6 is r3 after xor
  - Anything older than xor should read p3
  - Anything younger than xor should read p6 (until another insn writes r3)
- At commit of xor, no older instructions exist

# Commit Example

---


xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Commit Example

---

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Commit Example

---

add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4

Free-list

# Commit Example

---

sub r5 - r2 → r3  
addi r3 + 1 → r1

sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6

Free-list

# Commit Example

---

addi r3 + 1 → r1

addi p8 + 1 → p9

[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6
p1

Free-list

# Commit Example

---

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6
p1

Free-list

# Memory instructions

---



# Out-of-order is for loads!

---

Memory instructions are slow and hard for compilers to optimize!

- Cache hierarchy → Variable & long latency
- Memory disambiguation is hard → compilers can't optimize well

Out-of-order can **speculate to next load**

- Compilers cannot do this!
  - See *Decoupling Load for Nano-Instruction Set Computers*, Huang et al, ISCA'16

Out-of-order also handles variable latency well

- More on this later...

# Loads and stores: Implementation

---

Can't issue stores until instruction commits (*why not?*)

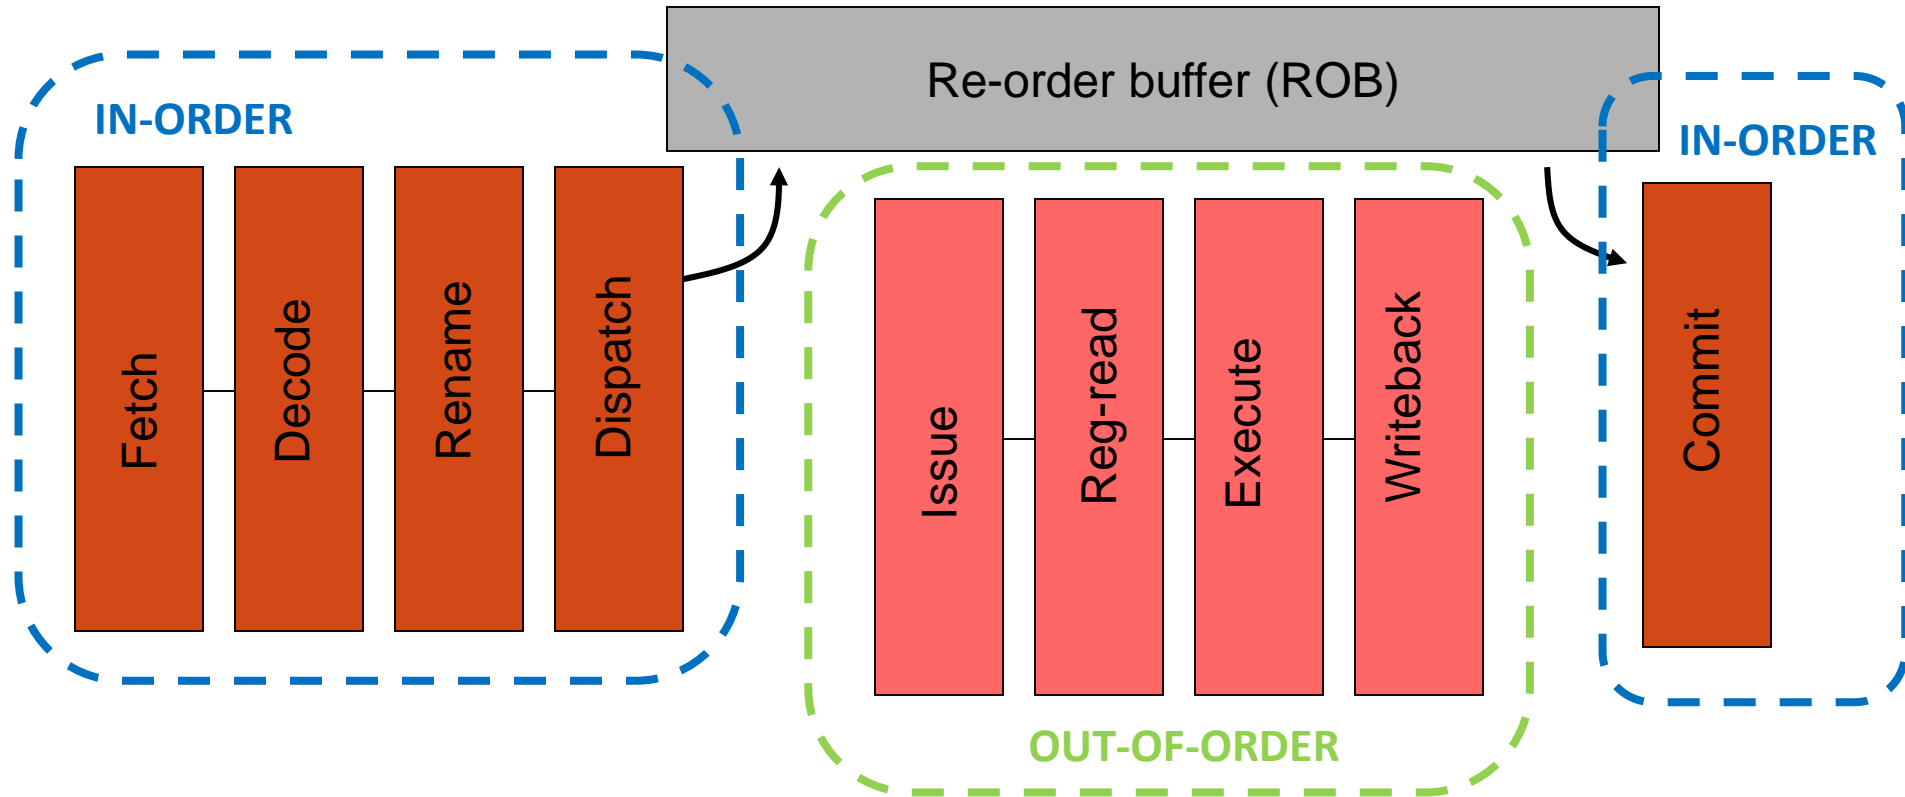
Load-store queue (LSQ) buffers executing stores

Store execution: reserve space in LSQ

Following loads must check LSQ, get **youngest earlier** value

- Complex associative lookup

# Out-of-order Pipeline



# OOO: Benefits & Challenges

---

# OOO Operation (Recap)

---

Fetch many instructions into instruction window

- Use branch prediction to [speculate](#) past (multiple) branches
- Flush pipeline on branch misprediction

Rename registers to [avoid false dependencies](#)

Execute instructions as soon as possible

- Register dependencies are known
- [Handling memory dependencies](#) is harder but doable (vs compiler)

“Commit” instructions [in order](#)

- Anything strange happens before commit, just flush the pipeline

# Out of Order: Benefits

---

OOO is very good at hiding **short** latencies

- Big IPC gain vs in-order superscalar

## Speculation

- Pull long-latency operations (loads, divides) across branches
- Gets better code schedule than compilers can achieve (often even in principle!)
  - Exponential growth in possible code paths, missing ISA mechanisms to split operations (active research area)

## Dynamism

- Change code schedule in response to variable latency ops

## Which matters more?

- Speculation gives >80% of the benefit [McFarlin et al, ASPLOS'13]

# OoO compared to other microarch

---

Scheduling done in hardware!

- Compilers must re-schedule code for different microarch in in-order

OoO takes the good ideas in dataflow and solves many of its challenges

- Limits dataflow to narrow window of related instructions → better locality (but less parallelism)
- Also called “restricted dataflow”

OoO runs “bad code” well – *it just works*

- Code compiled for old machine
- Pointer-chasing code hard for compilers to optimize

# Challenges for Out-of-Order Cores

---

## Design complexity

- More complicated than in-order? Certainly!
- But, we have managed to overcome the design complexity

## Clock frequency

- Can we build a “high ILP” machine at high clock frequency?
- Yep, with some additional pipe stages, clever design

## Limits to (efficiently) scaling the window and ILP

- Large physical register file
- Fast register renaming/wakeup/select/load queue/store queue
  - Lots of research ideas here
- Branch & memory dependency prediction limit effective window size
  - 95% branch mis-prediction: 1 in 20 branches, or 1 in 100 insns.
- Plus all the issues of building “wide” in-order superscalar

## Power efficiency

- Today, even mobile phone chips are out-of-order cores



# A Brief Survey of MicroArch Complexity

---

# I480 Pipeline

---

## Fetch

- Load 16-bytes of instruction into prefetch buffer

## Decode1

- Determine instruction length, instruction type

## Decode2

- Compute memory address
- Generate immediate operands

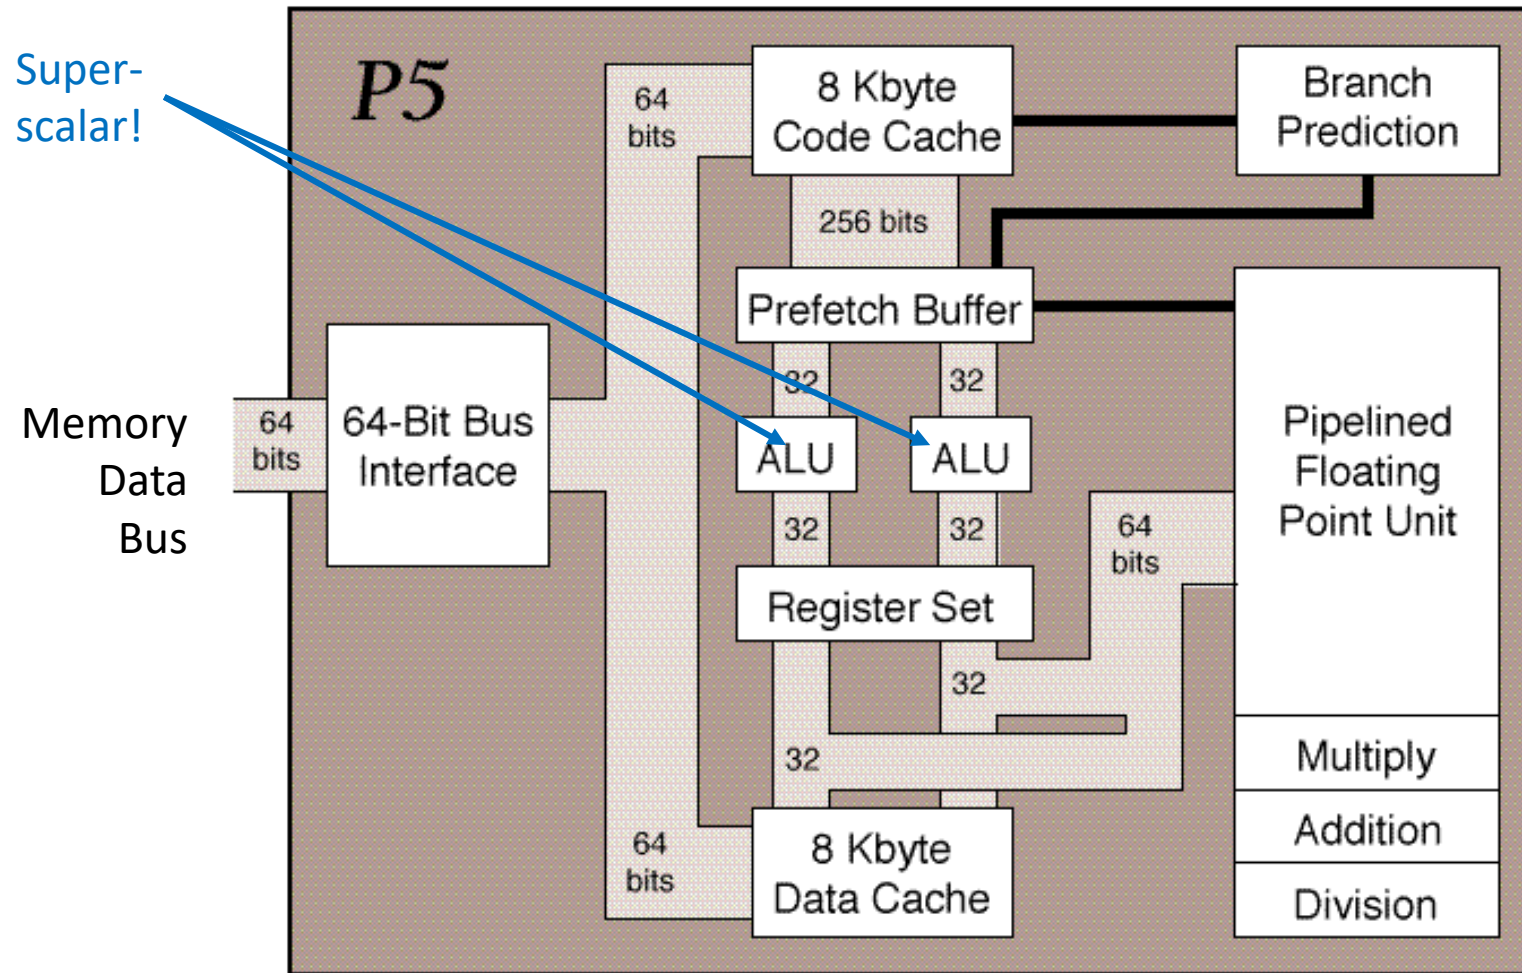
## Execute

- Register Read
- ALU operation
- Memory read/write

## Write-Back

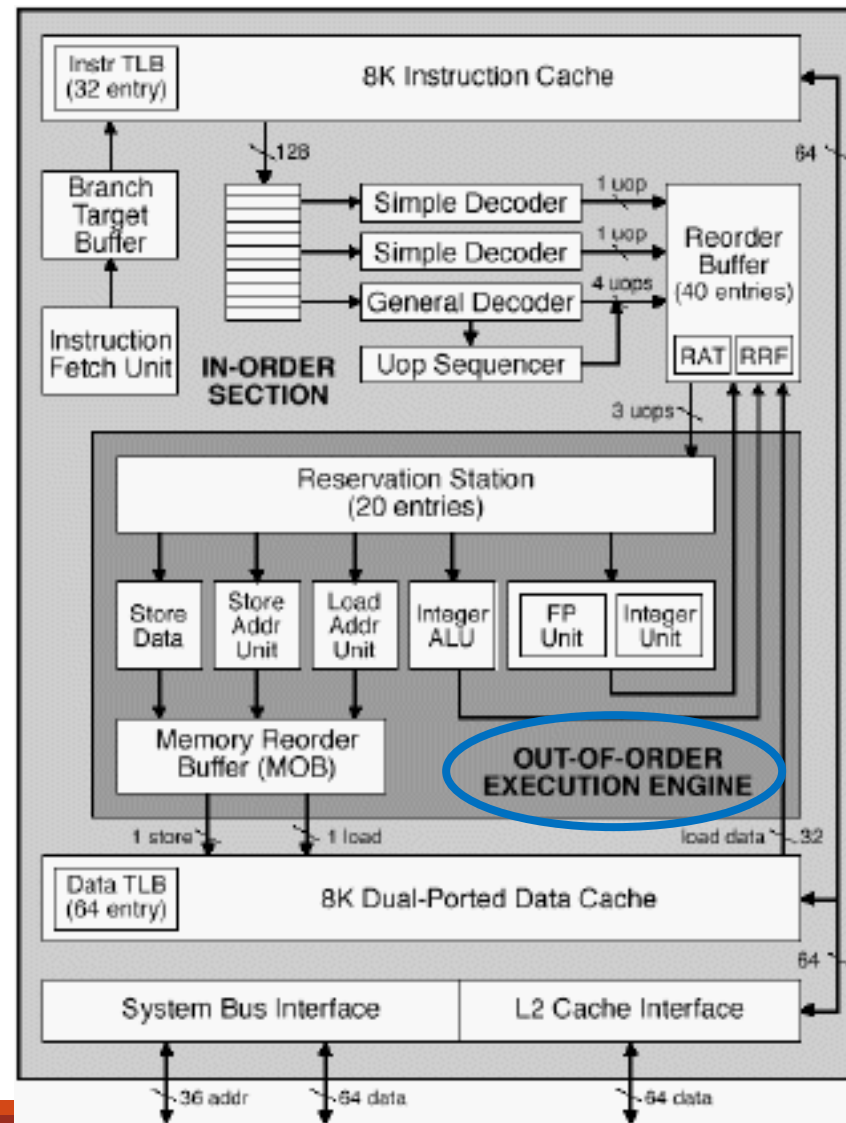
- Update register file

# Pentium Block Diagram




(Microprocessor Report 10/28/92)

# PentiumPro Block Diagram




Microprocessor Report  
2/16/95

# Core i7 Pipeline: Big Picture



(Images courtesy of Michael E. Thomadakis,  
Texas A&M Supercomputing Facility)

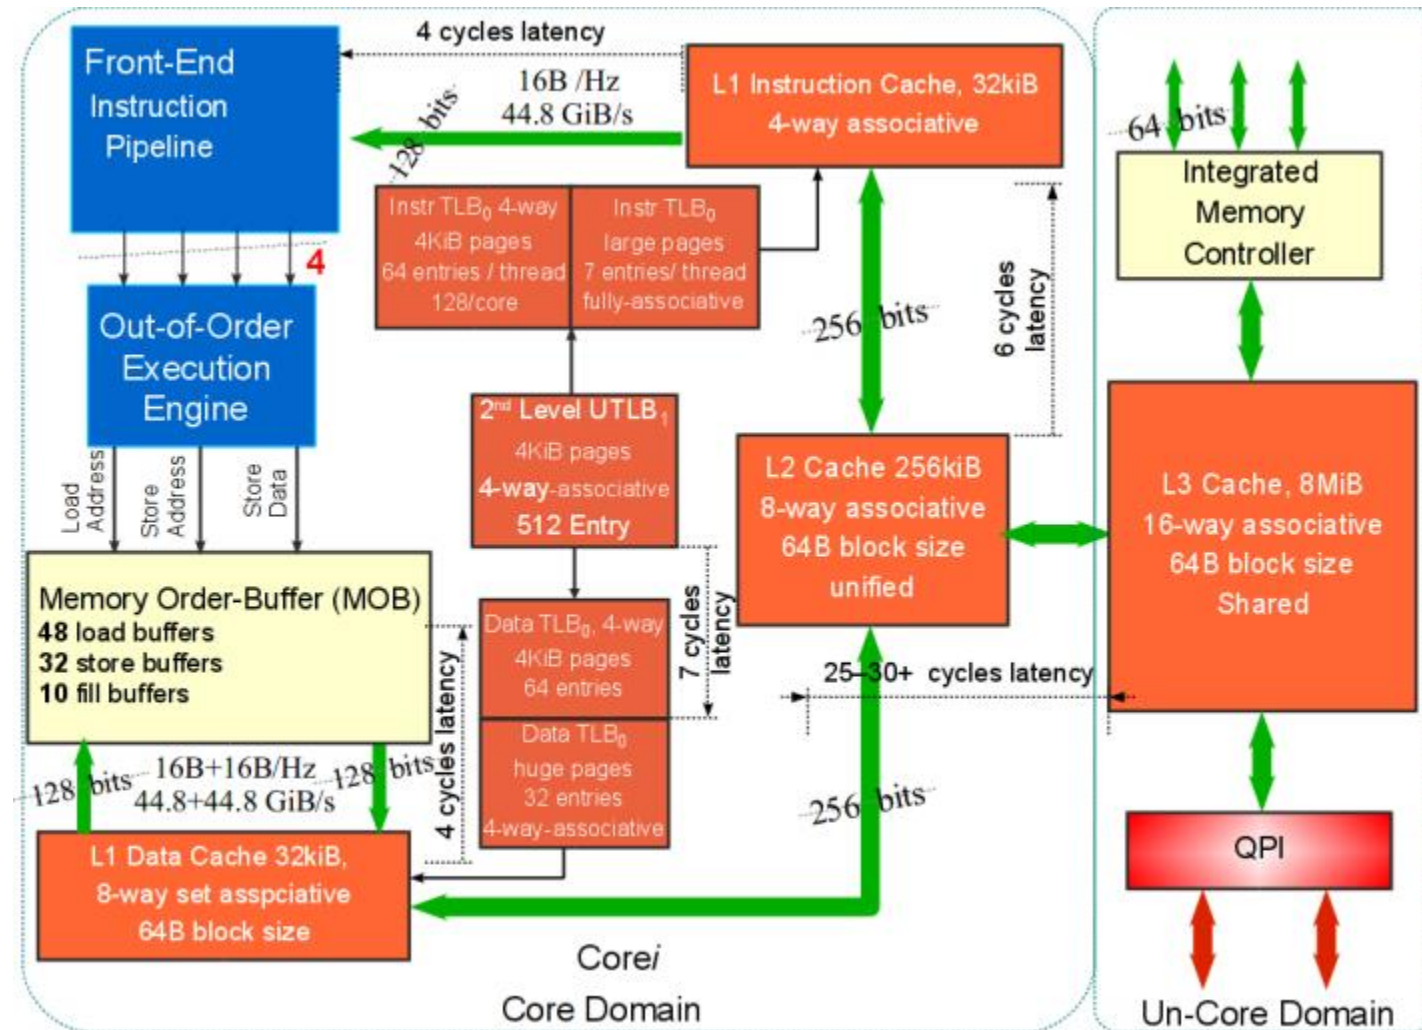
# Core i7 Pipeline: Front End





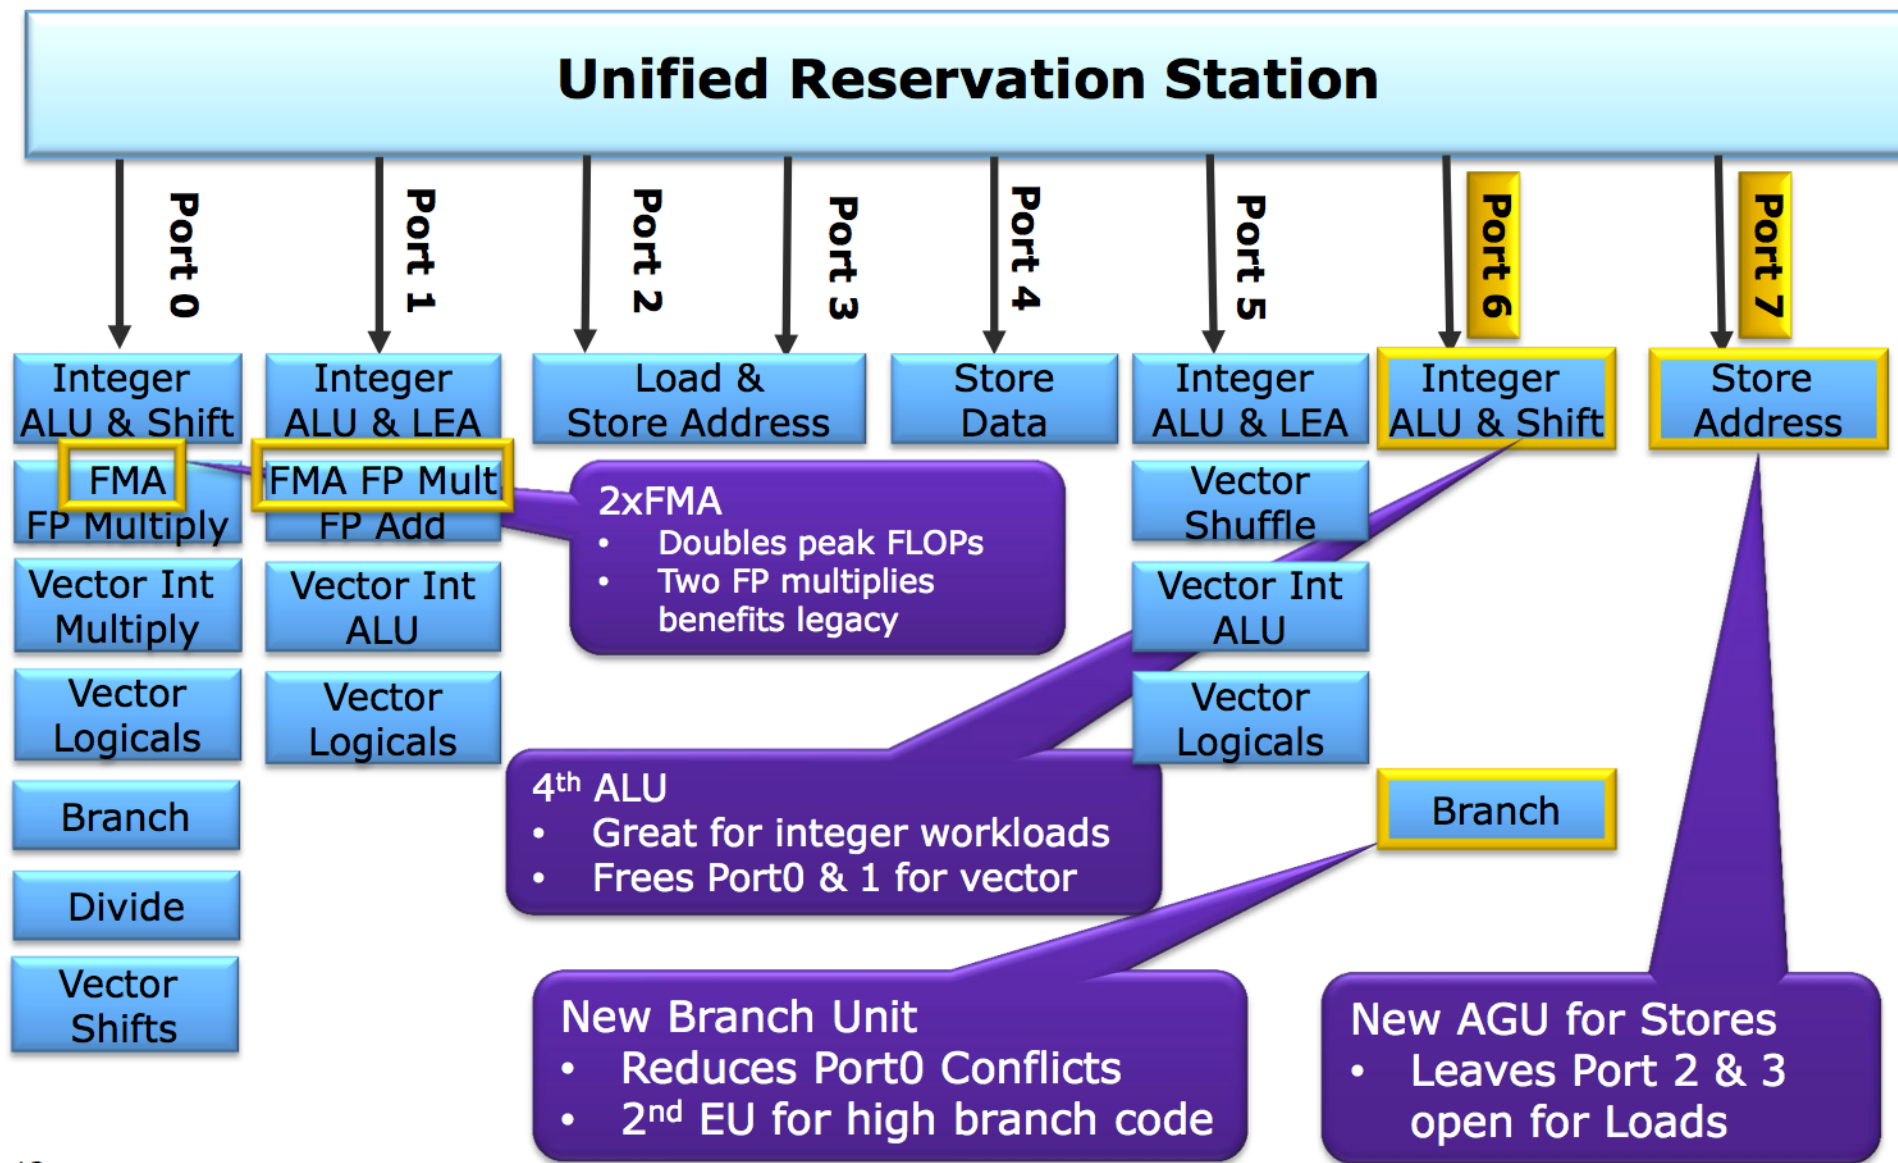


# Core i7 Pipeline: Memory Hierarchy












# Haswell Execution Unit Overview



# Haswell Buffer Sizes

Extract more parallelism in every generation

	Nehalem	Sandy Bridge	Haswell	
Out-of-order Window	128	168	192	
In-flight Loads	48	64	72	
In-flight Stores	32	36	42	
Scheduler Entries	36	54	60	
Integer Register File	N/A	160	168	
FP Register File	N/A	144	168	
Allocation Queue	28/thread	28/thread	56	

# OoO Execution is all around us

---

Qualcomm Krait processor (in phones)

- based on ARM Cortex A15 processor
- **out-of-order** 1.5GHz dual-core
- 3-wide fetch/decode
- 4-wide issue
- 11-stage integer pipeline
- 28nm process technology

# Living with Expensive Branches

---

## Mispredicted Branch Carries a High Cost

- Must flush many in-flight instructions
- Start fetching at correct target
- Will get worse with deeper and wider pipelines

## Impact on Programmer / Compiler

- Avoid conditionals when possible
  - Bit manipulation tricks
- Use special conditional-move instructions
  - Recent additions to many instruction sets
- Make branches predictable
  - Very low overhead when predicted correctly

# Branch Prediction Example

```
static void loop1() {
    int i;
    data_t abs_sum =
(data_t) 0;
    data_t prod = (data_t)
1;
    for (i = 0; i < CNT;
i++) {
        data_t x = dat[i];
        data_t ax;
        ax = ABS(x);
        abs_sum += ax;
        prod *= x;
    }
    answer = abs_sum+prod;
}
```

- Compute sum of absolute values
- Compute product of original values

```
#define ABS(x)  x < 0 ? -x :
x
```

## MIPS Code

```
0x6c4: 8c620000    lw  r2,0(r3)
0x6c8: 24840001    addiu r4,r4,1
0x6cc: 04410002    bgez
                    r2,0x6d8
0x6d0: 00a20018    mult    r5,r2
0x6d4: 00021023    subu
                    r2,r0,r2
0x6d8: 00002812    mflo    r5
0x6dc: 00c23021    addu
                    r6,r6,r2
0x6e0: 28820400    slti
                    r2,r4,1024
```

# Some Interesting Patterns

---

PPPPPPPP

+1 ...

- Should give perfect prediction

RRRRRRRR

-1 -1 +1 +1 +1 +1 -1 +1 -1 -1 +1 +1 -1 -1 +1 +1 +1 +1 +1 -1 -1 -1 +1 -1 ...

- Will mispredict 1/2 of the time

N\*N[PNPN]

-1 -1 -1 -1 -1 -1 -1 -1 +1 -1 +1 -1 +1 -1 +1 -1 +1 -1 +1 -1 +1 -1 +1 -1 ...

- Should alternate between states No! and No?

N\*P[PNPN]

-1 -1 -1 -1 -1 -1 -1 +1 +1 -1 +1 -1 +1 -1 +1 -1 +1 -1 +1 -1 +1 -1 +1 -1 ...

- Should alternate between states No? and Yes?

N\*N[PPNN]

-1 -1 -1 -1 -1 -1 -1 -1 +1 +1 -1 -1 +1 +1 -1 -1 +1 +1 -1 -1 +1 +1 -1 -1 ...

N\*P[PPNN]

-1 -1 -1 -1 -1 -1 -1 +1 +1 +1 -1 -1 +1 +1 -1 -1 +1 +1 -1 -1 +1 +1 -1 -1 ...

# Loop Performance (FP)

	R3000		PPC 604		Pentium	
Pattern	Cycles	Penalty	Cycles	Penalty	Cycles	Penalty
PPPPPPPPP	13.6	0	9.2	0	21.1	0
RRRRRRRRR	13.6	0	12.6	3.4	22.9	1.8
N*N [PNPN]	13.6	0	15.8	6.6	23.3	2.2
N*P [PNPN]	13.3	-0.3	15.9	6.7	24.3	3.2
N*N [PPNN]	13.3	-0.3	12.5	3.3	23.9	2.8
N*P [PPNN]	13.6	0	12.5	3.3	24.7	3.6

## Observations

- 604 has prediction rates 0%, 50%, and 100%
  - Expected 50% from N\*N[PNPN]
  - Expected 25% from N\*N[PPNN]
  - Loop so tight that speculate through single branch twice?
- Pentium appears to be more variable, ranging 0 to 100%

## Special Patterns Can be Worse than Random

- Only 50% of all people are “above average”

# Loop 1 Surprises

	R10000		Pentium II	
Pattern	Cycles	Penalty	Cycles	Penalty
PPPPPPPPP	3.5	0	11.9	0
RRRRRRRRR	3.5	0	19	7.1
N*N [PNPN]	3.5	0	12.5	0.6
N*P [PNPN]	3.5	0	13	1.1
N*N [PPNN]	3.5	0	12.4	0.5
N*P [PPNN]	3.5	0	12.2	0.3

## Pentium II

- Random shows clear penalty
- But others do well
  - More clever prediction algorithm

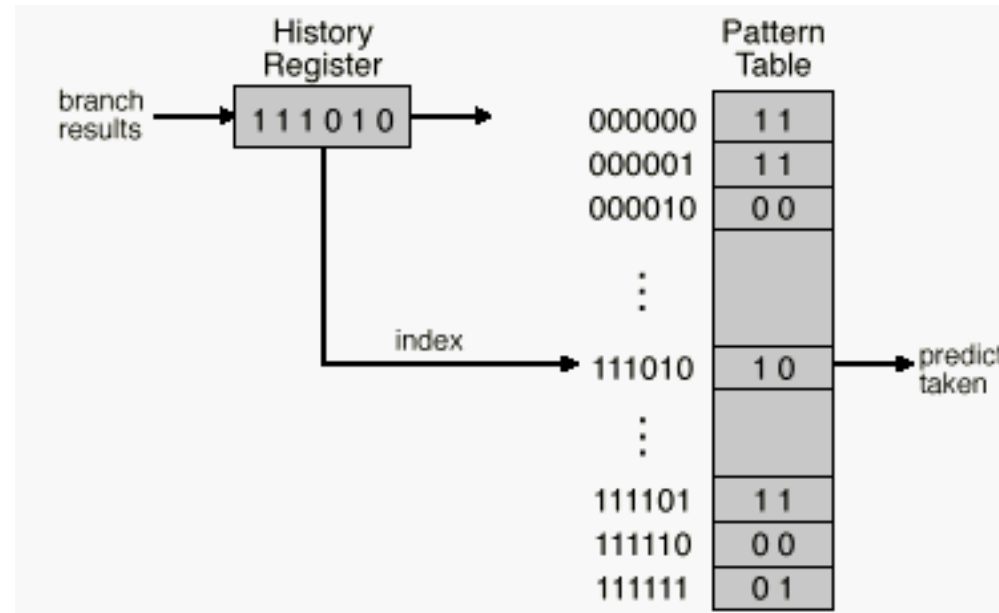
## R10000

- Has special “conditional move” instructions
- Compiler translates  $a = \text{Cond} ? \text{Texpr} : \text{Fexpr}$  into

```
a = Fexpr
temp = Texpr
CMOV(a, temp, Cond)
```
- Only valid if *Texpr* & *Fexpr* can't cause error



# P6 Branch Prediction




Microprocessor Report  
March 27, 1995

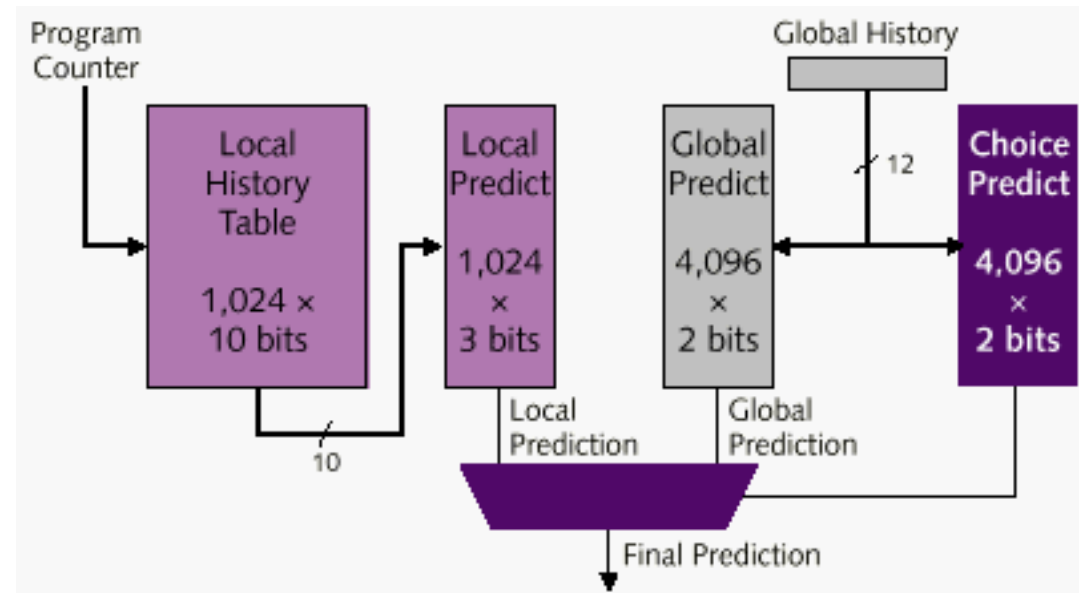
## Two-Level Scheme

- Yeh & Patt, ISCA '93
- Keep shift register showing past  $k$  outcomes for branch
- Use to index  $2^k$  entry table
- Each entry provides 2-bit, saturating counter predictor
- Very effective for any deterministic branching pattern

# Branch Prediction Comparisons

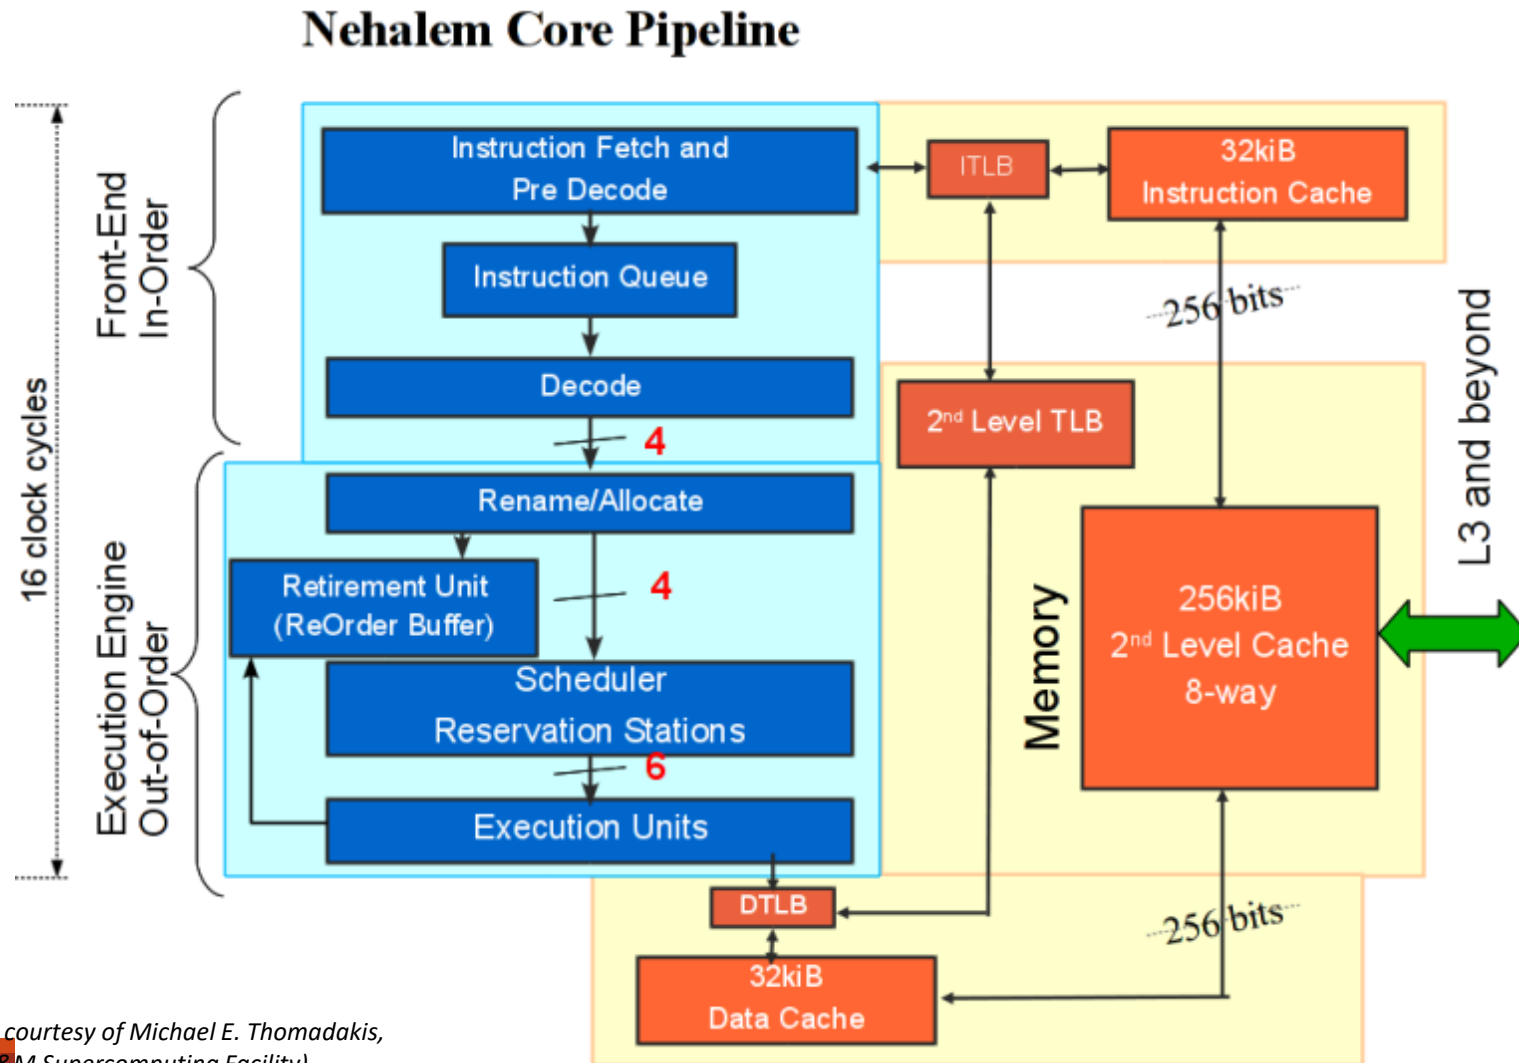


# 21264 Branch Prediction Logic



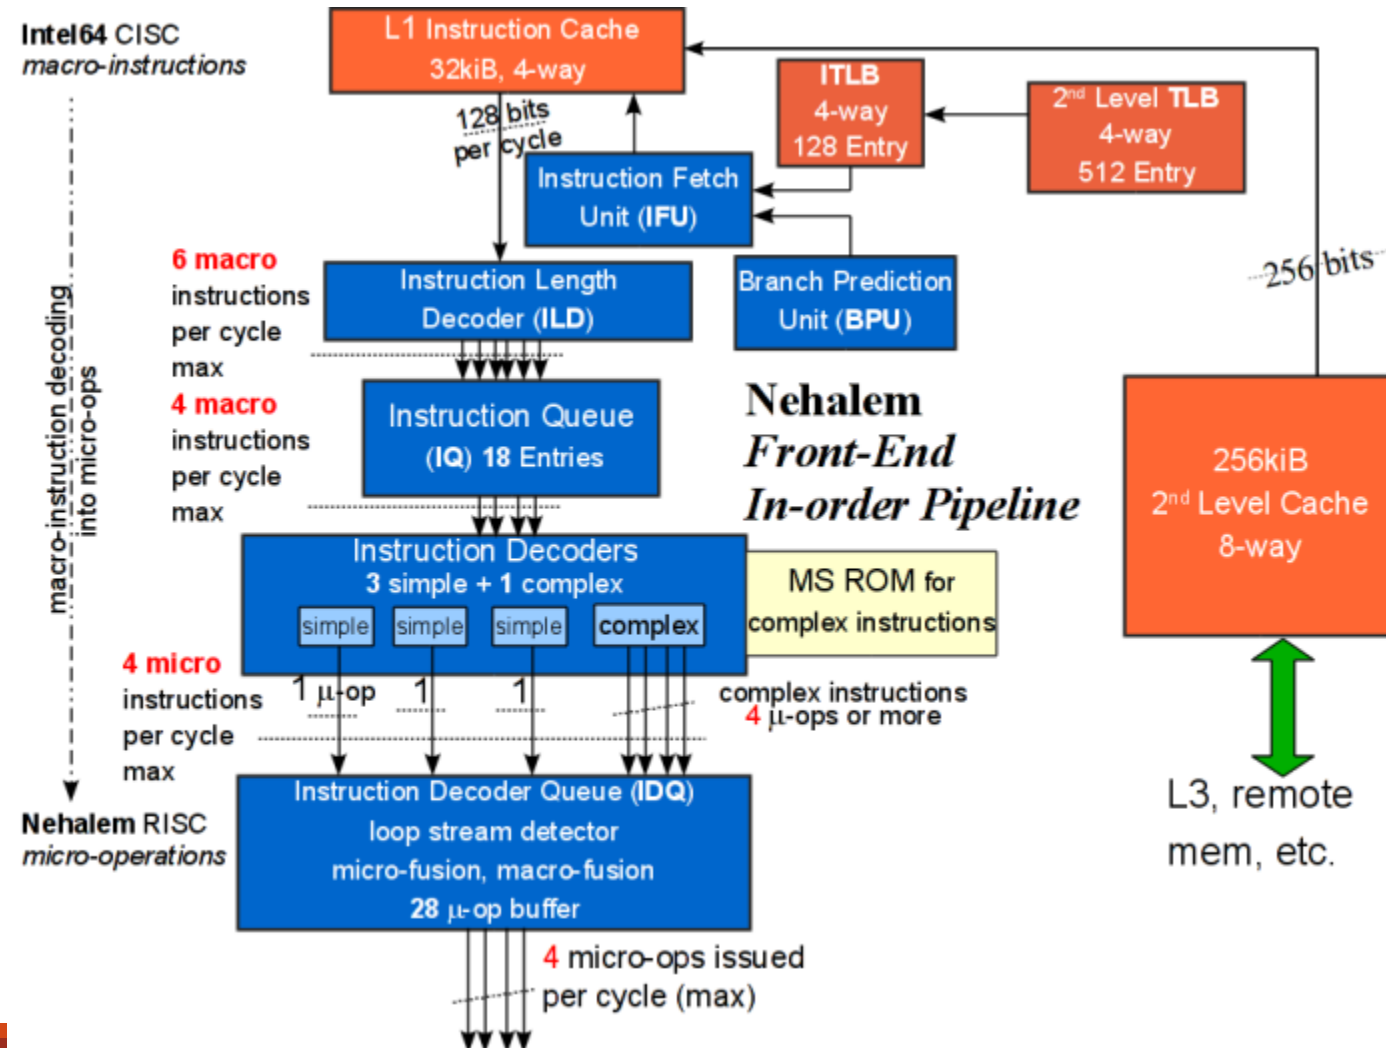
- Purpose: Predict whether or not branch taken
- 35Kb of prediction information
- 2% of total die size
- Claim 0.7--1.0% misprediction

# Core i7 Pipeline: Big Picture




(Images courtesy of Michael E. Thomadakis,  
Texas A&M Supercomputing Facility)

# Core i7 Pipeline: Front End



# Core i7 Pipeline: Execution Unit



# Core i7 Pipeline: Memory Hierarchy

