

# Computer Architecture

---

## Modern Out-of-Order Processors

Prof. Scott Beamer  
[sbeamer@ucsc.edu](mailto:sbeamer@ucsc.edu)

CSE 220  
<https://canvas.ucsc.edu/courses/56561>

**Baskin  
Engineering**  
UC SANTA CRUZ



# Increasing Performance via ILP

---

- By Little's Law, to improve instruction throughput, if we can't improve latency, we will need to increase *parallelism*
- Types of parallelism
  - *Temporal* – multiple things in flight on same hardware – pipelining
  - *Spatial* – multiple hardware units – superscalar
- Challenge to executing more instructions at once is thus:
  - Identifying which instructions can execute in parallel
  - Providing instructions with the needed data
- ILP enhancement techniques in this course
  - *Pipelining* – including forwarding
  - *Caches* – fewer memory stalls means more instructions in flight
  - *Dynamic scheduling* – including scoreboarding (today) and Tomasulo
  - *Branch prediction* – speculate branches to uncover more instructions
  - *Memory prefetching* – further reduce memory stalls
  - *TLP* – use multiple threads/cores to execute more

# Types of Data Hazards

Consider executing a sequence of

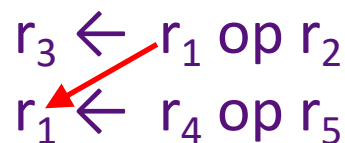
$$r_k \leftarrow r_i \text{ op } r_j$$

type of instructions

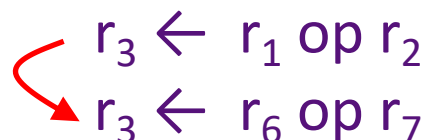
Data-dependence

$$\begin{array}{ll} r_3 \leftarrow r_1 \text{ op } r_2 & \text{Read-after-Write} \\ r_5 \leftarrow r_3 \text{ op } r_4 & \text{(RAW) hazard} \end{array}$$


Anti-dependence

$$\begin{array}{ll} r_3 \leftarrow r_1 \text{ op } r_2 & \text{Write-after-Read} \\ r_1 \leftarrow r_4 \text{ op } r_5 & \text{(WAR) hazard} \end{array}$$


Output-dependence

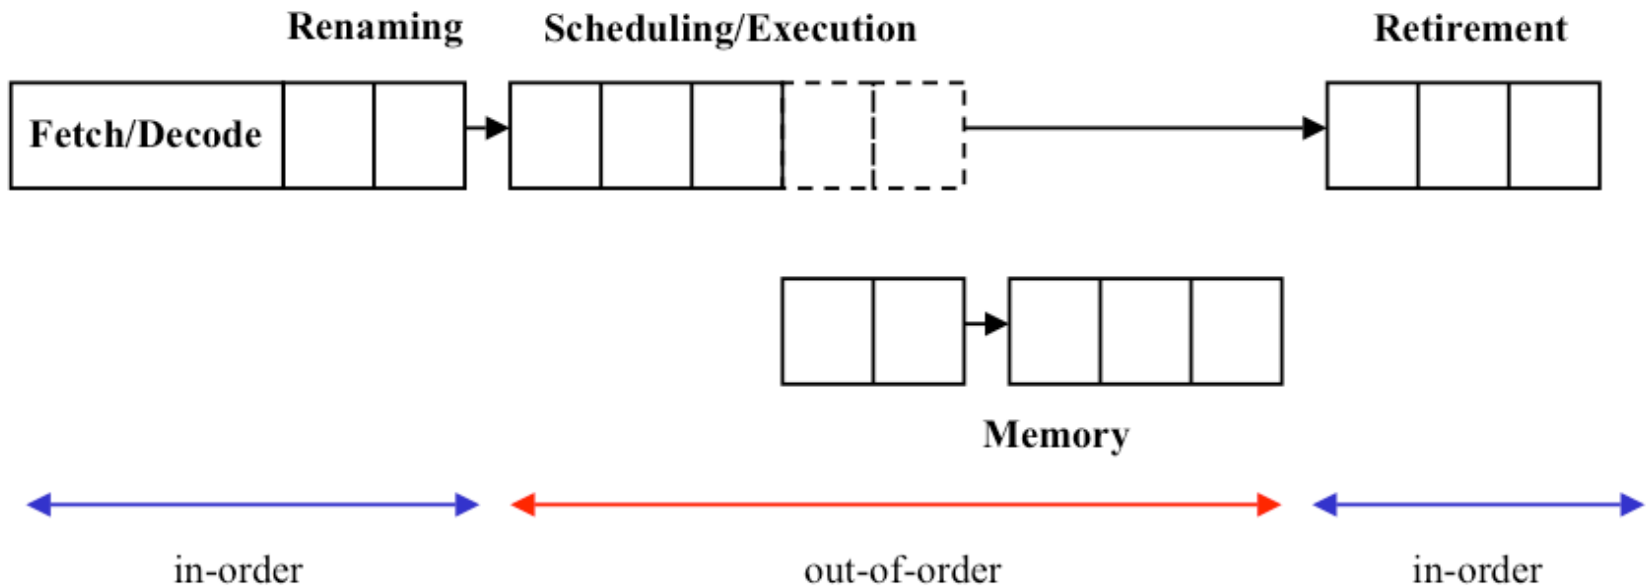
$$\begin{array}{ll} r_3 \leftarrow r_1 \text{ op } r_2 & \text{Write-after-Write} \\ r_3 \leftarrow r_6 \text{ op } r_7 & \text{(WAW) hazard} \end{array}$$


# Roadmap for Exploiting ILP

---

- Single-cycle CPU
  - Directly implements abstraction of in order & 1 instruction at a time
- Pipelined (e.g. 5-stage RISC)
  - Can overlap execution of multiple instructions, but everything in order
  - Forwarding helps with some RAW, in-order eliminates WAR/WAW
- Scoreboard
  - Fetch/Decode/Issue in-order, but allow out-of-order exec. & completion
  - Helpful for accommodating long latency functional units
  - Stalls for any hazard (RAW,WAR,WAW), but instr. can pass another
- Tomasulo
  - Uses register renaming to eliminate WAR & WAW
- Modern Out-of-Order (OoO)
  - Supports precise-exceptions
  - Simplifies/streamlines/centralizes OoO structures
  - Only limited by RAW hazards

# Out-of-Order Processor Flow



# Issue Window / Reservation Stations

---

- Hold instructions until executed
  - Instructions wait for operands and free functional unit
- Lifecycle of instruction in Issue Window
  - Inserted into window after rename (same time as into ROB)
  - *Wakeup* – operands become available (marked ready)
  - *Select* – chosen from ready instructions (and FU free) to execute
  - *Issue* – instruction was selected, sent to FU and removed
- Wakeup and select often combined into same pipeline stage
- Issue window is typically much smaller than ROB
  - ROB holds instructions until committed
  - Issue window holds instructions until start execution (issue)
  - Instructions remaining in issue window may not be contiguous

# Modern OoO Example

## Reorder Buffer (ROB)

```
ld    R1, 0(R2)
addi  R2, R2, 8
sub   R3, R3, R1
bne   R2, R0, L
ld    R1, 0(R2)
```

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>						
I <sub>2</sub>						
I <sub>3</sub>						
I <sub>4</sub>						
I <sub>5</sub>						

### Frontend RAT

R0	
R1	
R2	
R3	

### Issue Window

Unit	Src.1	Src. 1 Ready	Src. 2	Src. 2 Ready

**Free List:**

### Retirement RAT

R0	
R1	
R2	
R3	

# Modern OoO Example

## Reorder Buffer (ROB)

```

ld    R1, 0(R2)
addi  R2, R2, 8
sub   R3, R3, R1
bne   R2, R0, L
ld    R1, 0(R2)
    
```

H →

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>	ld	R1	P2	P3	i0	
I <sub>2</sub>						
I <sub>3</sub>						
I <sub>4</sub>						
I <sub>5</sub>						

## Frontend RAT

R0	P0
R1	<del>P6</del> P2
R2	P3
R3	P5

## Issue Window

Unit	Src.1	Src. 1 Ready	Src. 2	Src. 2 Ready
MEM	P3	✓	i0	✓

Free List: ~~P2~~, P1, P7, P8, P9

## Retirement RAT

R0	P0
R1	P6
R2	P3
R3	P5



# Modern OoO Example

## Reorder Buffer (ROB)

~~ld R1, 0(R2)~~  
 addi R2, R2, 8  
 sub R3, R3, R1  
 bne R2, R0, L  
 ld R1, 0(R2)

T →

H →

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>	ld	R1	P2	P3	i0	
I <sub>2</sub>	addi	R2	P1	P3	i8	
I <sub>3</sub>						
I <sub>4</sub>						
I <sub>5</sub>						

## Frontend RAT

R0	P0
R1	<del>P6</del> P2
R2	<del>P3</del> P1
R3	P5

## Issue Window

Unit	Src.1	Src. 1 Ready	Src. 2	Src. 2 Ready
ALU	P3	✓	i8	✓

Free List: ~~P2~~, ~~P1~~, P7, P8, P9

## Retirement RAT

R0	P0
R1	P6
R2	P3
R3	P5

# Modern OoO Example

## Reorder Buffer (ROB)

~~ld~~ R1, 0(R2)  
 addi R2, R2, 8  
 sub R3, R3, R1  
 bne R2, R0, L  
 ld R1, 0(R2)

T →

H →

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I1	ld	R1	P2	P3	i0	
I2	addi	R2	P1	P3	i8	
I3	sub	R3	P7	P5	P2	
I4						
I5						

### Frontend RAT

R0	P0
R1	<del>P6</del> P2
R2	<del>P3</del> P1
R3	<del>P5</del> P7

### Issue Window

Unit	Src. 1	Src. 1 Ready	Src. 2	Src. 2 Ready
ALU	P5	✓	P2	

Free List: ~~P2~~, ~~P1~~, ~~P7~~, P8, P9

### Retirement RAT

R0	P0
R1	P6
R2	P3
R3	P5

# Modern OoO Example

## Reorder Buffer (ROB)

~~ld~~ R1, 0(R2)  
 addi R2, R2, 8  
 sub R3, R3, R1  
 bne R2, R0, L  
 ld R1, 0(R2)

T →

H →

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>	ld	R1	P2	P3	i0	
I <sub>2</sub>	addi	R2	P1	P3	i8	✓
I <sub>3</sub>	sub	R3	P7	P5	P2	
I <sub>4</sub>	bne	-	-	P1	P0	
I <sub>5</sub>						

### Frontend RAT

R0	P0
R1	<del>P6</del> P2
R2	<del>P3</del> P1
R3	<del>P5</del> P7

### Issue Window

Unit	Src.1	Src. 1 Ready	Src. 2	Src. 2 Ready
BRA	P1	✓	P0	✓
ALU	P5	✓	P2	

Free List: ~~P2~~, ~~P1~~, ~~P7~~, P8, P9

### Retirement RAT

R0	P0
R1	P6
R2	P3
R3	P5

# Modern OoO Example

## Reorder Buffer (ROB)

~~ld~~ R1, 0(R2)  
 addi R2, R2, 8  
 sub R3, R3, R1  
 bne R2, R0, L  
 ld R1, 0(R2)

T →  
 H →

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>	ld	R1	P2	P3	i0	✓
I <sub>2</sub>	addi	R2	P1	P3	i8	✓
I <sub>3</sub>	sub	R3	P7	P5	P2	
I <sub>4</sub>	bne	-	-	P1	P0	
I <sub>5</sub>	ld	R1	P8	P1	i0	

### Frontend RAT

R0	P0
R1	<del>P6</del> P2 P8
R2	<del>P3</del> P1
R3	<del>P5</del> P7

### Issue Window

Unit	Src.1	Src. 1 Ready	Src. 2	Src. 2 Ready
MEM	P1	✓	i0	✓
ALU	P5	✓	P2	✓

Free List: ~~P2~~, ~~P1~~, ~~P7~~, ~~P8~~, P9, P6

### Retirement RAT

R0	P0
R1	<del>P6</del> P2
R2	P3
R3	P5

# Modern OoO Example

## Reorder Buffer (ROB)

ld R1, 0(R2)  
 addi R2, R2, 8  
 sub R3, R3, R1  
 bne R2, R0, L  
 ld R1, 0(R2)

H → I<sub>1</sub>  
 T → I<sub>2</sub>

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>						
I <sub>2</sub>	addi	R2	P1	P3	i8	✓
I <sub>3</sub>	sub	R3	P7	P5	P2	
I <sub>4</sub>	bne	-	-	P1	P0	✓
I <sub>5</sub>	ld	R1	P8	P1	i0	

### Frontend RAT

R0	P0
R1	<del>P6</del> P2 P8
R2	<del>P3</del> P1
R3	<del>P5</del> P7

### Issue Window

Unit	Src. 1	Src. 1 Ready	Src. 2	Src. 2 Ready

Free List: ~~P2~~, ~~P1~~, ~~P7~~, ~~P8~~, P9, P6, P3

### Retirement RAT

R0	P0
R1	<del>P6</del> P2
R2	<del>P3</del> P1
R3	P5

# Modern OoO Example

## Reorder Buffer (ROB)

L: ld R1, 0(R2)  
 addi R2, R2, 8  
 sub R3, R3, R1  
 bne R2, R0, L  
 ld R1, 0(R2)

H → I<sub>1</sub>

T → I<sub>3</sub>

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>						
I <sub>2</sub>						
I <sub>3</sub>	sub	R3	P7	P5	P2	
I <sub>4</sub>	bne	-	-	P1	P0	✓
I <sub>5</sub>	ld	R1	P8	P1	P0	

### Frontend RAT

R0	P0
R1	<del>P6</del> P2 P8
R2	<del>P3</del> P1
R3	<del>P5</del> P7

### Issue Window

Unit	Src. 1	Src. 1 Ready	Src. 2	Src. 2 Ready

Free List: ~~P2~~, ~~P1~~, ~~P7~~, ~~P8~~, P9, P6, P3

### Retirement RAT

R0	P0
R1	<del>P6</del> P2
R2	<del>P3</del> P1
R3	P5

# Modern OoO Example

## Reorder Buffer (ROB)

ld R1, 0(R2)  
 addi R2, R2, 8  
 sub R3, R3, R1  
 bne R2, R0, L  
 ld R1, 0(R2)

H → I<sub>1</sub>

T → I<sub>3</sub>

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>						
I <sub>2</sub>						
I <sub>3</sub>	sub	R3	P7	P5	P2	✓
I <sub>4</sub>	bne	-	-	P1	P0	✓
I <sub>5</sub>	ld	R1	P8	P1	i0	

### Frontend RAT

R0	P0
R1	<del>P6</del> P2 P8
R2	<del>P3</del> P1
R3	<del>P5</del> P7

### Issue Window

Unit	Src. 1	Src. 1 Ready	Src. 2	Src. 2 Ready

Free List: ~~P2~~, ~~P1~~, ~~P7~~, ~~P8~~, P9, P6, P3

### Retirement RAT

R0	P0
R1	<del>P6</del> P2
R2	<del>P3</del> P1
R3	P5

# Modern OoO Example

## Reorder Buffer (ROB)

L: ld R1, 0(R2)  
 addi R2, R2, 8  
 sub R3, R3, R1  
 bne R2, R0, L  
 ld R1, 0(R2)

H → I<sub>1</sub>  
 T → I<sub>4</sub>

	Op.	Logical Dest.	Physical Dest.	Physical Src. 1	Physical Src. 2	Executed
I <sub>1</sub>						
I <sub>2</sub>						
I <sub>3</sub>						
I <sub>4</sub>	bne	-	-	P1	P0	✓
I <sub>5</sub>	ld	R1	P8	P1	i0	

## Frontend RAT

R0	P0
R1	<del>P6</del> P2 P8
R2	<del>P3</del> P1
R3	<del>P5</del> P7

## Issue Window

Unit	Src. 1 Ready	Src. 2 Ready

Free List: ~~P2, P1, P7, P8~~, P9, P6, P3, P5

## Retirement RAT

R0	P0
R1	<del>P6</del> P2
R2	<del>P3</del> P1
R3	<del>P5</del> P7



# Exception Handling

---

- Issue Logic executes out-of-order, ROB brings original order back
- *Key: no store visible outside the CPU until retirement/commit*
  - all previous instructions committed
    - architectural state up to date
  - this and subsequent instructions not committed
    - architectural state not compromised
  - false exceptions from mispredicted path never emerge
- Once excepting instruction makes it to commit stage:
  - flush all uncommitted instructions
  - save precise architectural state
  - handle exception
  - restore architectural state and initiate fetch

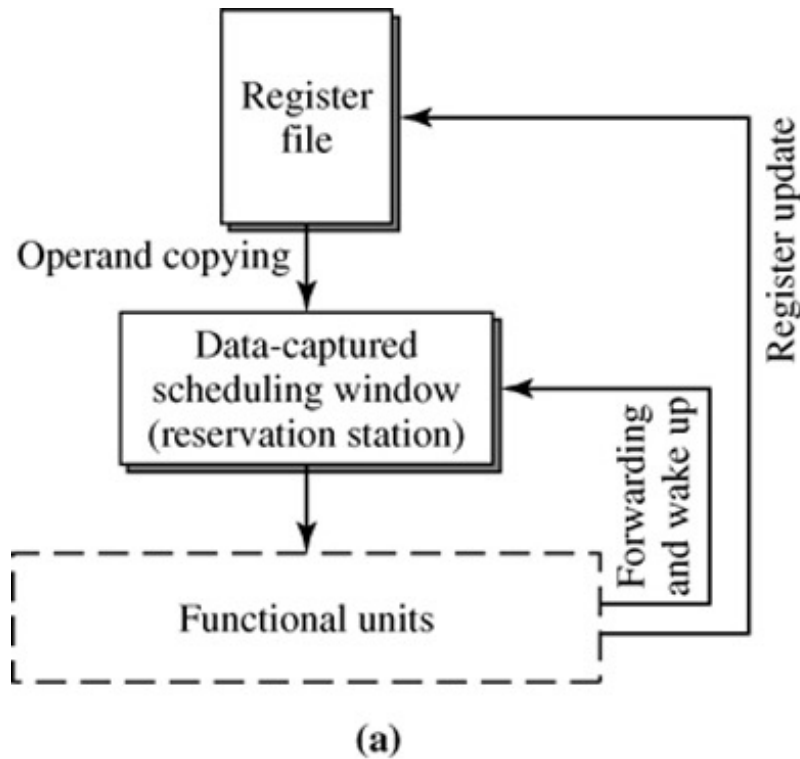
# Peer Instruction Question

---

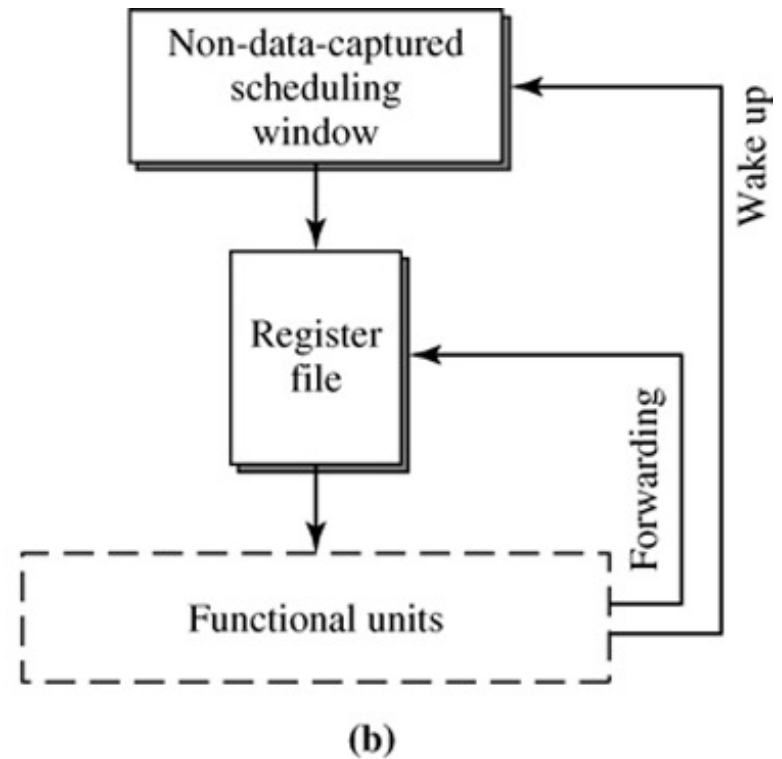
To count how many instructions in flight write to the same architectural register, you should check:

- A) Count values that match in Frontend RAT
- B) Count values that match in Retirement RAT
- C) Count # entries in ROB with matching logical destination
- D) Sum of all of above

# When to Read the Register File?



Data in ROB



Unified Physical Register File

# Timeline Samples

## Modern OoO, no FWD

a=b+c	IF	RN	WS	R0	R1	EX	WB					
d=a+3		IF	RN	--	--	--	WS	R0	R1	EX	WB	

## Modern OoO with FWD

a=b+c	IF	RN	WS	R0	R1	EX	WB					
d=a+3		IF	RN	WS	R0	R1	EX	WB				

## Modern OoO with FWD

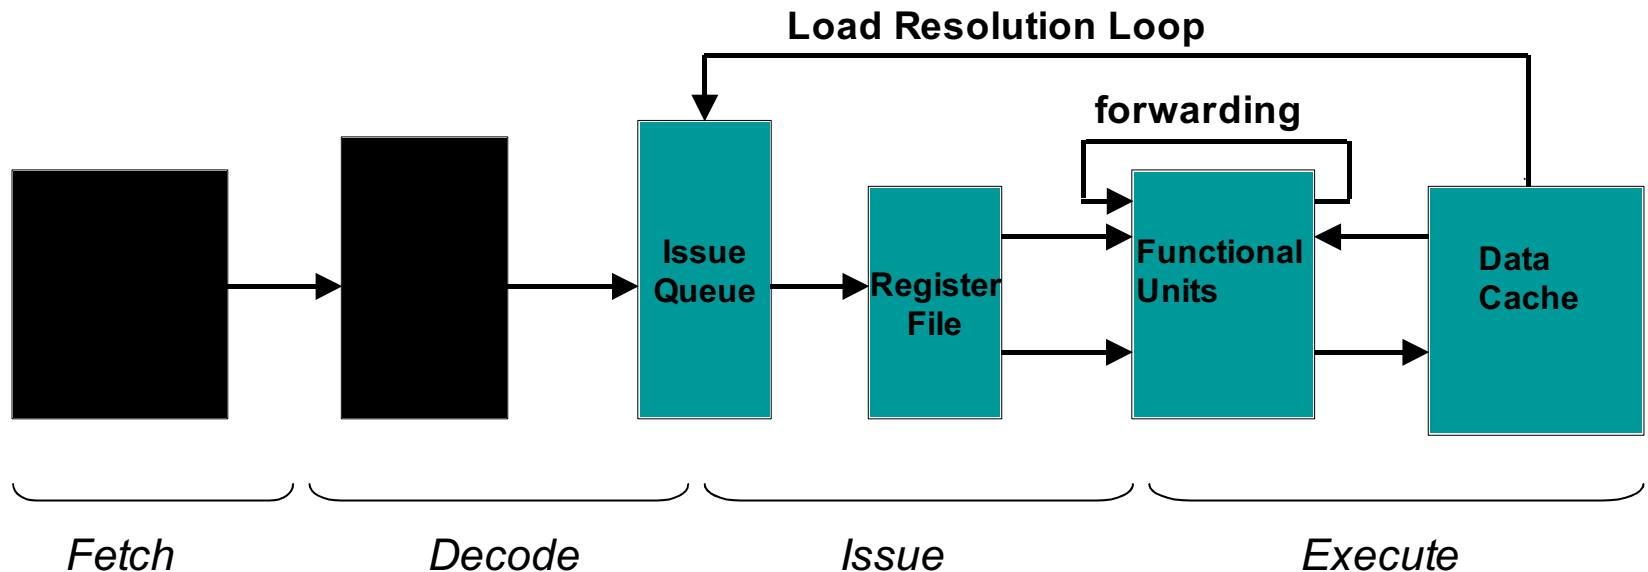
a=a*b	IF	RN	WS	R0	R1	E0	E1	WB				
d=a+3		IF	RN	--	WS	R0	R1	EX	WB			

## Modern OoO with FWD (no load hit speculation)

a=ld(f)	IF	RN	WS	R0	R1	M0	M1	WB				
d=a+3		IF	RN	--	--	--	--	WS	R0	R1	EX	WB

# Load-Hit Speculation

- We can not wait to wakeup dependent instructions until load finish
- Must have feature for in-order and out-order processors
  - Becoming more important as the pipeline depth increases



# Timeline Samples

## Modern OoO, no FWD

a=b+c IF RN WS R0 R1 EX WB

d=a+3 IF RN RN RN RN WS R0 R1 EX WB

## Modern OoO with FWD

a=b+c IF RN WS R0 R1 EX WB

d=a+3 IF RN WS R0 R1 EX WB

## Modern OoO with FWD

a=a\*b IF RN WS R0 R1 E0 E1 WB

d=a+3 IF RN RN WS R0 R1 EX WB

## Modern OoO with FWD (no load hit speculation)

a=ld(f) IF RN WS R0 R1 M0 M1 WB

d=a+3 IF RN RN RN RN RN RN WS R0 R1 EX WB

## Modern OoO with FWD (load hit speculation with cache hit)

a=ld(f) IF RN WS R0 R1 M0 M1 WB

d=a+3 IF RN RN WS R0 R1 EX WB

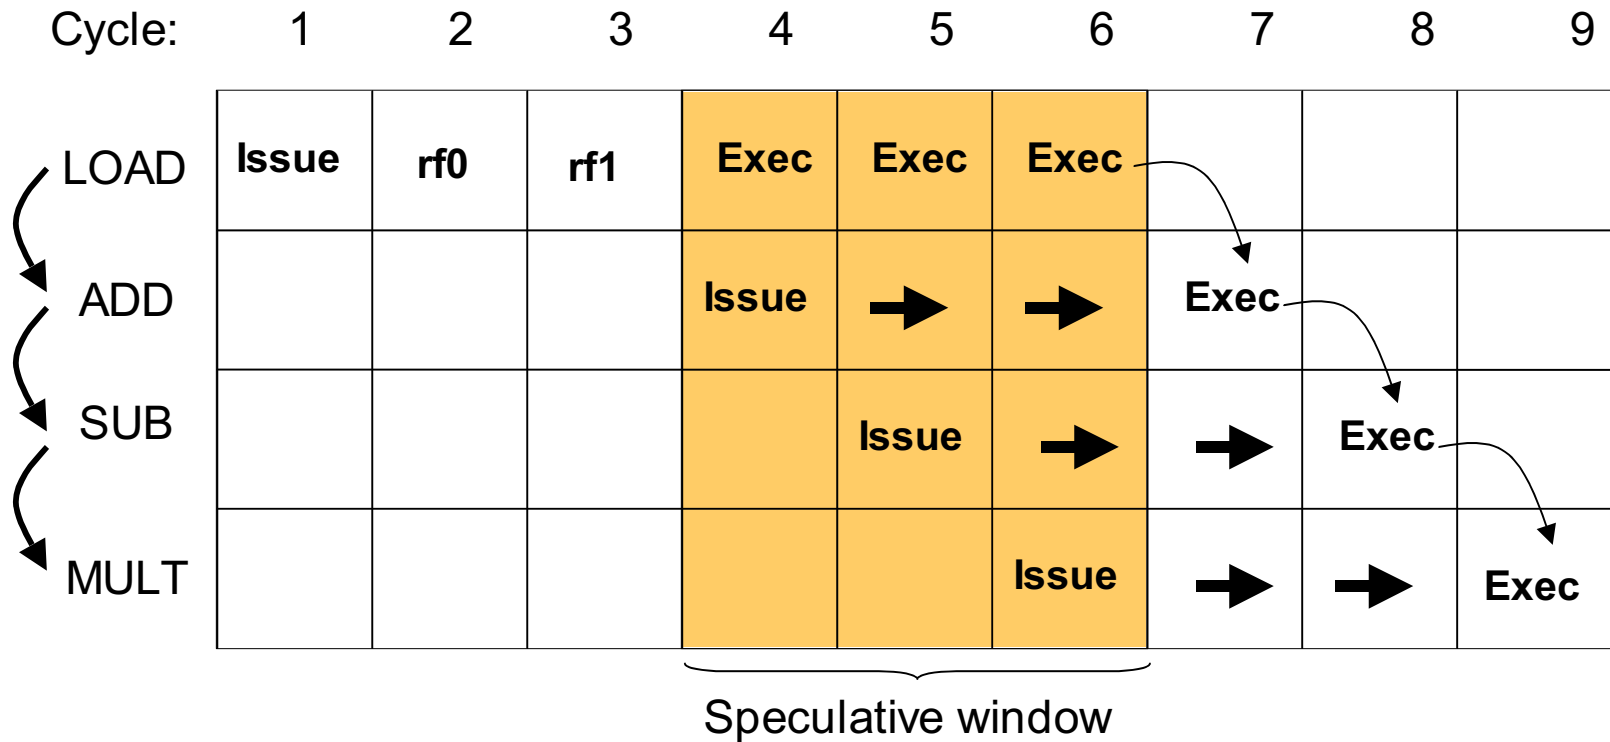
## Modern OoO with FWD (load hit speculation with cache miss)

a=ld(f) IF RN WS R0 R1 M0 M1 M2 M3 WB

d=a+3 IF RN RN WS R0 R1 EX WB

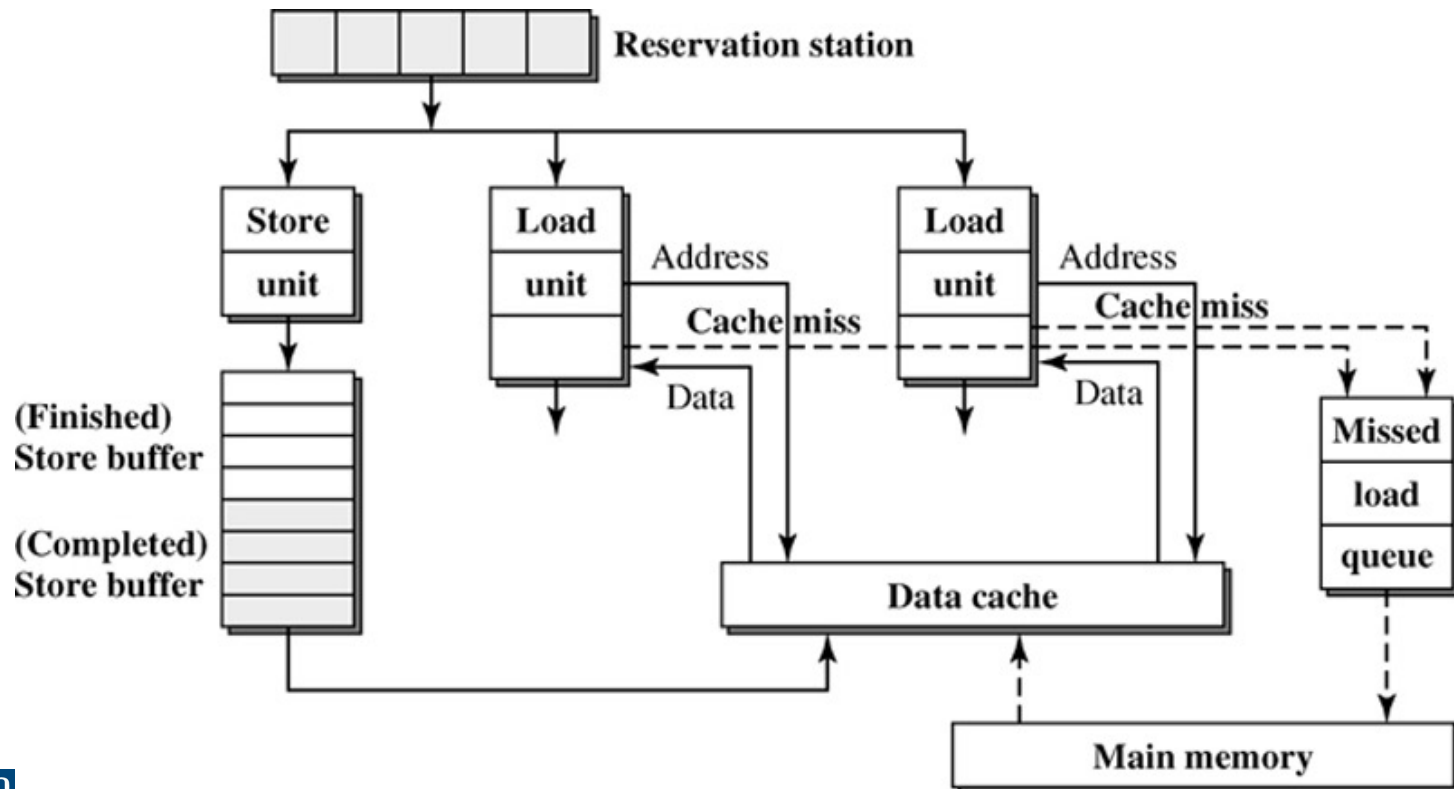
d=a+3 -- -- -- -- -- -- -- WS R0 R1 EX WB

# Load-Hit Speculation: Speculative Window



# Memory Speculation, Why should I care?


- Can you exec. memory instructions (loads & stores) out-of-order?
- Between 10-30% slowdown with no memory speculation
- All high performance OoO cores implement it





# No Memory Speculation Design

- Naïve - serialize all stores and loads
  - Exploits less memory-level parallelism
- Want to start loads sooner when possible
- When is it safe to load?
  - When confident no prior (older) store will write to same address (*alias*)
  - Thus, need to know not only load address, but also all prior store addresses
  - This is conservative OoO



```
1: sd r6, 0(r1)
2: ld r7, 0(r2)
3: ld r8, 0(r3)
4: sd r9, 0(r4)
5: sd r10, 0(r5)
```

# Load/Store Ordering

---

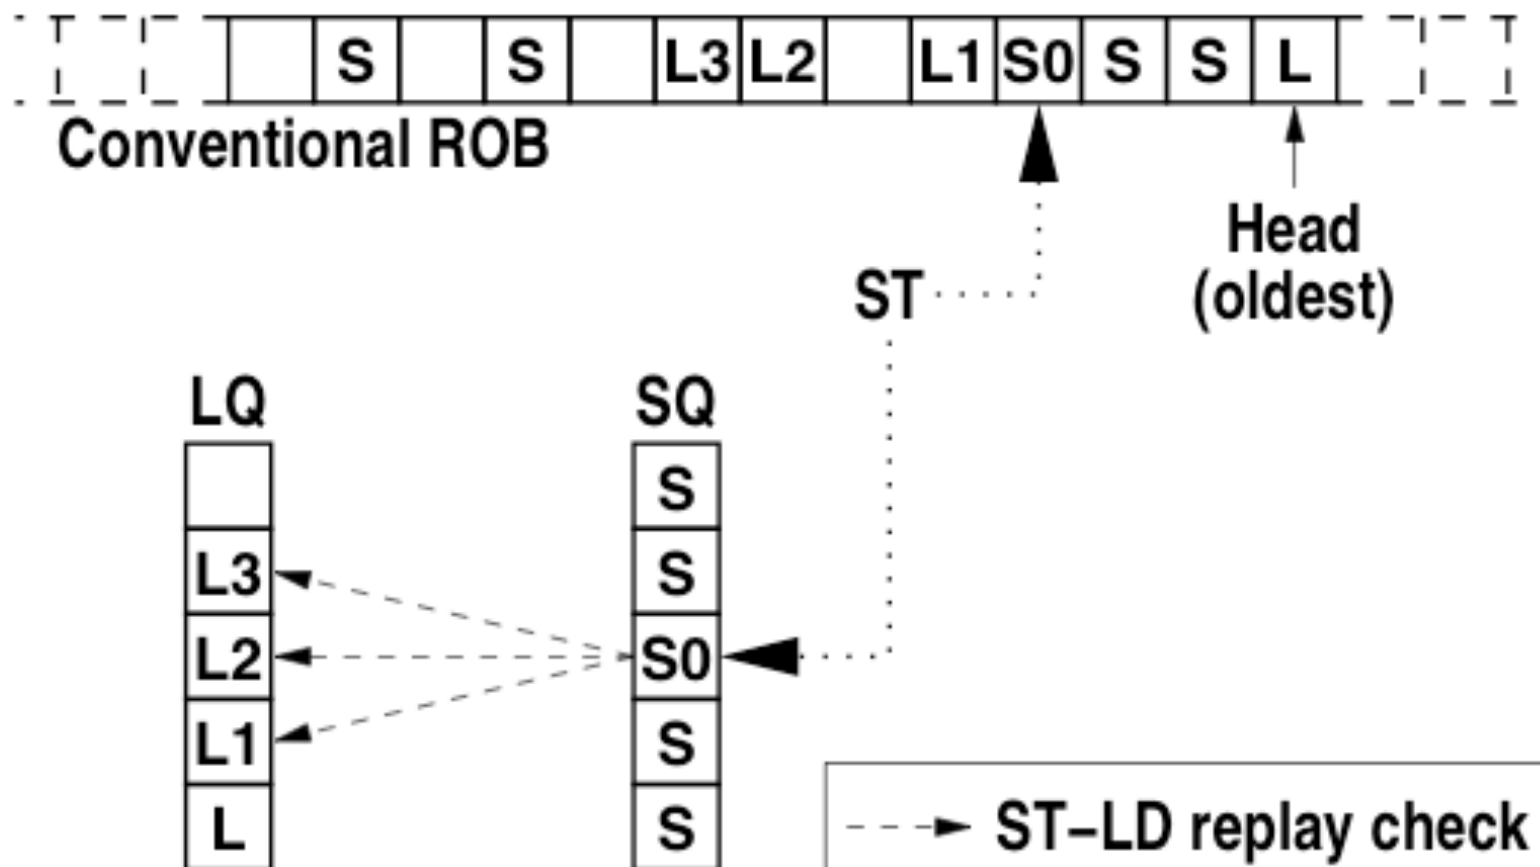
- Loads may execute out of order
  - starting loads early diminishes impact of miss penalty
- Stores always perform at commit
  - WAW – taken care of by in-order commit
- Must respect store-load ordering (ST-LD Replay)
  - RAW – must not execute (load) before store completes
    - could forward value from store buffer (not supported here)
  - WAR – automatically enforced by in-order stores at commit
- RAW enforcement: do not issue load if
  - earlier stores to overlapping address
  - *earlier stores to unresolved address*

# Memory Speculation

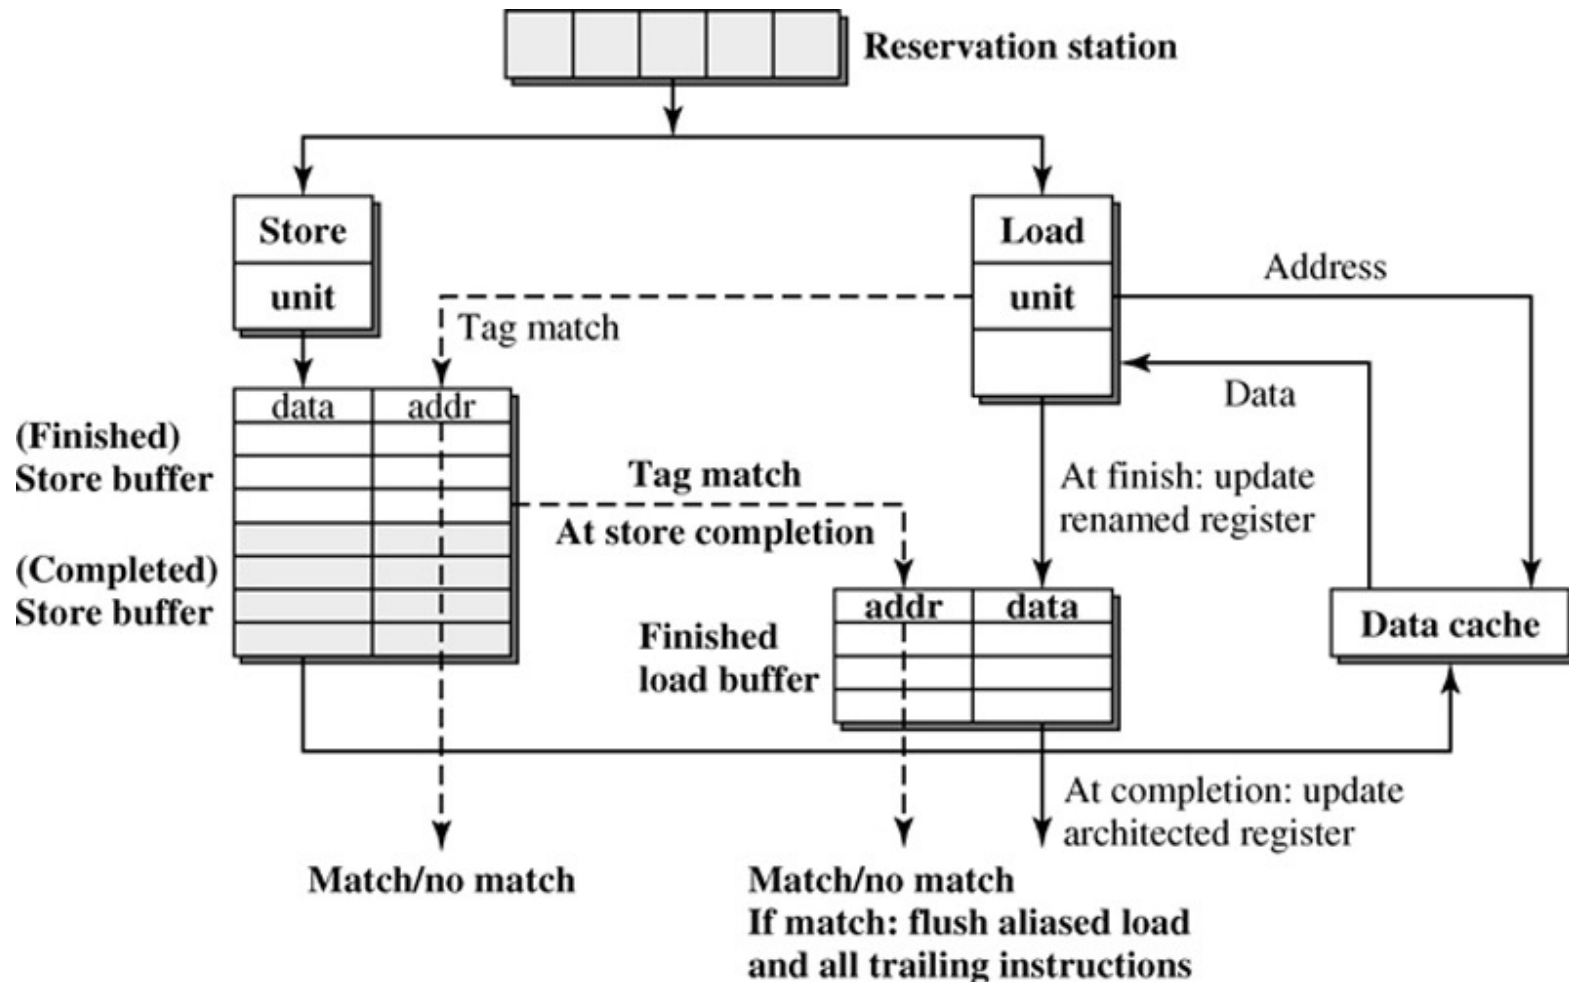
---

- If a prior store address is unknown, can *speculate* that the load will not have the same address (alias)
  - If speculation is right, got more things in flight (yay!)
  - If speculation is wrong, will need to recover
- To speculate, will need the following:
  - Load Queue (LQ)
    - Keeps all the loads addresses
  - Store Queue (SQ)
    - Keeps all the stores addresses
  - Store Data Queue (SDQ)
    - Queue associated with the store queue. Keeps the stores data

# ST-LD Replay Check



# MPD Diagram for LSQ



# Store Completion Buffer (SCB)

---

- Holds stores (data and address) after retirement
  - Once store is in SCB, is retired
- SCB is essentially a store buffer (as covered in cache lectures)
- Holding store in buffer helps...
  - Hide variable latency of cache, so don't want to slow retire
    - Could be a cache miss (on a system that is write-allocate & write-back)
  - Coalesce multiple stores to same address range
- Visibility of store
  - Pre-retirement – nobody outside processor can see store
    - Instructions within processor can read data via forwarding from SDQ
  - Post-retirement – “written” to memory
    - Requests for that data will check cache and SCB

# Memory Speculation Summary

---

- Allow loads to start early by *speculating* will not alias with prior stores to unknown addresses
- Load remains in LQ until committed
- When store commits, checks for alias in LQ
  - If no match, safe
  - If addresses match, perform recovery
    - Load should have read data from store since same address
    - If load was before store in program order, would already be out of LQ
- Remember: *commit stage is in order*
  - If load was earlier in program order (WAR), will commit first, and will be gone from LQ when store commits
  - If store was earlier in program order (RAW), the process above double checks load got right value (and flushes if wrong)
- Read *Modern Processor Design* 5.3.2 – 5.3.3 for more

# Acknowledgements

---

- These slides contain course material initially created by Jose Renau for CMPE 202